

# **An Investigation Into the Cognitive Processing of Magnitude in Data Visualisations**

A thesis submitted to The University of Manchester for the degree of  
Doctor of Philosophy  
in the Faculty of Biology, Medicine and Health

**2023**

**Duncan Bradley**

School of Health Sciences

# Table of Contents

<b>List of Figures</b>	<b>4</b>
<b>List of Tables</b>	<b>5</b>
<b>Abstract</b>	<b>6</b>
<b>Declaration</b>	<b>7</b>
<b>Copyright</b>	<b>8</b>
<b>Acknowledgements</b>	<b>9</b>
<b>1 Introduction</b>	<b>10</b>
1.1 Research Motivation and Objective . . . . .	11
1.2 Contributions . . . . .	11
1.3 Overview of the Thesis . . . . .	12
References . . . . .	13
<b>2 Literature Review</b>	<b>14</b>
2.1 A Brief History of Data Visualisation . . . . .	14
2.2 Data Visualisation Formats and the Grammar of Graphics . . . . .	18
2.3 Data Visualisation Software and the Influence of Default Settings . . . . .	19
2.4 Popular Guidance on Effective Data Visualisation Design . . . . .	21
2.5 Empirical Research on Data Visualisation . . . . .	24
2.6 Perceptual Precision in Data Visualisations . . . . .	26
2.7 Beyond Perceptual Precision . . . . .	26
2.8 Manipulating Axes in Data Visualisations . . . . .	29
2.9 Misleading Data Visualisations . . . . .	32
2.10 Data Visualisation Literacy . . . . .	32
2.11 Interpreting Absolute Magnitude . . . . .	33
2.12 The Present Thesis . . . . .	36
References . . . . .	36
<b>3 Experimental, Analytical, and Computational Methodology</b>	<b>47</b>
3.1 Introduction . . . . .	47
3.2 Experimental Methodology . . . . .	47
3.2.1 Experimental Design and Ecological Validity . . . . .	48
3.2.2 Sample Size and Generalisability . . . . .	48
3.3 Analytical Methodology . . . . .	49
3.3.1 Linear Mixed-Effects Models . . . . .	49
3.3.2 Fixed Effects and Random Effects . . . . .	49
3.3.3 Benefits of Mixed-Effects Models . . . . .	50

3.3.4	Approaches to Model Construction . . . . .	50
3.3.5	An Automated Approach . . . . .	50
3.3.6	Reporting Analyses . . . . .	51
3.4	Reproducibility . . . . .	51
3.5	Sharing Code and Data . . . . .	52
3.5.1	The Importance of Public Sharing . . . . .	54
3.6	Effective Programming Practices . . . . .	55
3.6.1	Literate Programming and Dynamic Documents . . . . .	55
3.7	Computational Environments . . . . .	56
3.7.1	Capturing Computational Environments Using Containers . . . . .	57
3.7.2	Rocker for Capturing R Environments . . . . .	58
3.7.3	Comparing Containers with Virtual Machines . . . . .	58
3.8	Pragmatism Over Perfectionism . . . . .	59
3.9	The Approach to Reproducibility in This Thesis . . . . .	59
3.9.1	Data, Code, and Dynamic Documents . . . . .	59
3.9.2	Docker Containers . . . . .	60
3.9.3	Experiment Resources . . . . .	61
3.10	Conclusion . . . . .	62
	References . . . . .	62
	<b>Preface to Chapter 4</b>	<b>69</b>
	References . . . . .	70

Word Count:

## List of Figures

2.1	John Snow's Map of Cholera Deaths (1855) . . . . .	15
2.2	Florence Nightingale's Rose Diagram (1857) . . . . .	16
2.3	Charles Joseph Minard's Map of Napoleon's Russian Campaign (1861) . .	16
2.4	Warming Stripes by Ed Hawkins, showing average global temperature data from 1850 to 2018. . . . .	17
2.5	Eight different ways of displaying the same dataset, to illustrate the Gram- mar of Graphics. . . . .	20
2.6	A reproduction of a data visualisation used by Tufte (1986) to illustrate the 'lie-factor', originally published in the Los Angeles Times. . . . .	23
2.7	Visualisations used to demonstrate greater perceptual averaging ability for colour encoding (low spatial frequency), than position encoding (high spatial frequency). . . . .	27
2.8	Examples of data points positioned 'inside' and 'outside' a bar showing an average value, and the equivalent plots with confidence intervals. . . . .	28
2.9	Examples of truncated and non-truncated line and bar charts, showing the same data. . . . .	29
2.10	Stimuli used to demonstrate that including additional shaded bars (grey) reduces interpretations of the magnitude of other values (blue), compared to blank space. . . . .	35

## List of Tables

## Abstract

Data visualisations are effective communication tools which leverage the power of the human visual and cognitive systems. However, data visualisation design choices can substantially influence interpretations of presented information. Interpreting data visualisations does not just involve accurately comprehending values, but often also involves drawing inferences about data and making subjective judgements. Therefore, developing an understanding of cognitive processing of data visualisations is important for guiding the construction of effective and faithful representations of data.

Data visualisations can convey many different features of a dataset. One such feature is the absolute magnitude of plotted values: how large or small they are. This thesis investigates cognitive mechanisms involved in judging absolute magnitude in data visualisations, revealing the effects of design choices on interpretation.

The first set of experiments in this thesis (three experiments) explores the role of axis limits in informing magnitude judgements. Manipulating axis limits in dot plots causes the same data points to appear near the top or bottom of the visualisation. Participants' responses revealed an association between higher positions and higher magnitude ratings, indicating a bias in interpretation. Two further experiments employing dot plots with inverted y-axes indicated that impressions of magnitude were driven primarily by the relative positions of data points within axis limits, not their absolute physical positions.

The second experiment in this thesis extends inquiry into the role of axis limits to choropleth maps. Manipulating the limits of accompanying colour legends alters the framing of presented data, without changing how plotted values appeared. Extending the colour legend's upper limit beyond the maximum plotted value resulted in lower magnitude ratings. This demonstrates that interpretations of absolute magnitude are informed by surrounding context, not just by the appearance of plotted values.

The final set of experiments in this thesis (two experiments) explores how additional knowledge about plotted data can inform interpretations of magnitude. Denominators provide numerical context relevant to magnitude judgements. Extending a bar chart's axis beyond plotted data to incorporate a denominator value elicited lower magnitude ratings, compared to bar charts' default settings. Omitting denominator information from accompanying text substantially increased this bias. This illustrates that additional knowledge about a dataset diminishes the roles of axis limits in informing impressions of magnitude.

This work was conducted with a focus on computational reproducibility. In addition to sharing data and analysis code, this involved facilitating the ability to reproduce the computational environment used for analysis. This approach, which increases openness and transparency in research, is also discussed in detail.

Through experimental research, this thesis reveals how the framing of values within axes informs judgements of their absolute magnitudes. The results provide insight into the cognitive processing of magnitude in data visualisations, wherein context shapes viewers' inferences. This illustrates how inevitable subjectivity in data visualisation design can influence a data visualisation's appearance and its message. Data visualisation designers should consider the graphical representation of absolute magnitude and, where appropriate, employ axes ranges which faithfully convey this aspect of data.

## Declaration

- One experiment reported in Chapter 4 was conducted in collaboration with Gabriel Strain, who submitted a report on this research written independently in partial fulfilment of his MRes Psychology degree at the University of Manchester (2021). The report written within this thesis is my own work.
- The experiment reported in Chapter 5 was conducted in collaboration with Boshuo Zhang, who submitted a report on this research written independently in partial fulfilment of his MSc Advanced Computer Science degree at the University of Manchester (2022). The report written within this thesis is my own work.
- One experiment reported in Chapter 6 was conducted in collaboration with Harvey Schneider, who submitted a report on this research written independently in partial fulfilment of his MRes Psychology degree at the University of Manchester (2022). The report written within this thesis is my own work.
- No other portion of the work referred to in the thesis has been submitted in support of an application for another degree or qualification of this or any other university or other institute of learning.

## Copyright

- i. The author of this thesis (including any appendices and/or schedules to this thesis) owns certain copyright or related rights in it (the “Copyright”) and they have given the University of Manchester certain rights to use such Copyright, including for administrative purposes.
- ii. Copies of this thesis, either in full or in extracts and whether in hard or electronic copy, may be made **only** in accordance with the Copyright, Designs and Patents Act 1988 (as amended) and regulations issued under it or, where appropriate, in accordance with licensing agreements which the University has from time to time. This page must form part of any such copies made.
- iii. The ownership of certain Copyright, patents, designs, trademarks and other intellectual property (the “Intellectual Property”) and any reproductions of copyright works in the thesis, for example graphs and tables (“Reproductions”), which may be described in this thesis, may not be owned by the author and may be owned by third parties. Such Intellectual Property and Reproductions cannot and must not be made available for use without the prior written permission of the owner(s) of the relevant Intellectual Property and/or Reproductions.
- iv. Further information on the conditions under which disclosure, publication and commercialisation of this thesis, the Copyright and any Intellectual Property and/or Reproductions described in it may take place is available in the University IP Policy (see <https://documents.manchester.ac.uk/DocuInfo.aspx?DocID=24420>), in any relevant Thesis restriction declarations deposited in the University Library, the University Library’s regulations (see <https://www.library.manchester.ac.uk/about/regulations/>) and in the University’s policy on Presentation of Theses.



## Acknowledgements

I would like to thank my supervisors, Andrew Stewart and Caroline Jay. Andrew's mentorship during my BSc and MRes degrees was foundational in motivating my PhD, and he has generously offered his time, expertise, and his fresh outlook, throughout the entire process. Caroline's perspective has also been extremely valuable and her wisdom has been instrumental in improving my work. I am very grateful to have been given the opportunity to collaborate with such a dedicated team. I am also indebted to Jen McBride for the considerate and pragmatic guidance which she offered as my advisor.

For productive conversations, practical support, and camaraderie, thanks to Libby Le-luan, Gabe Strain, and Paul Stott. Thanks also to all in 124e, and to Tristan Johnston-Wood and Sam Langton, for sharing experiences of the PhD process. I am grateful to Boshuo Zhang and Harvey Schneider for their hard work in helping to program experimental stimuli during their Master's programmes.

## Chapter 1

# Introduction

Data visualisations help people make sense of numbers. Whilst a list of numbers may allude to an upwards trend, presenting those numbers *as* visual phenomena can aid interpretation. A data visualisation can systematically depict the precise numerical values, facilitating mental processing of this information.

As *external* representations of data, visualisations reduce perceptual and cognitive burdens in interpretation (Scaife and Rogers, 1996). By imparting efficiency and clarity, data visualisations support pattern-recognition and reasoning. However, a single dataset can be depicted in numerous ways, and different designs can vary widely in their effectiveness (Franconeri et al., 2021). Thus, data visualisation's strength can also be its vulnerability. Outsourcing cognitive processes to a graphical depiction leaves a viewer at the mercy of the chosen method of visual representation. Thus, understanding successful design is crucial.

The effectiveness of data visualisations can be defined in many ways, encompassing their various objectives, which include informing, persuading, engaging, and promoting retention of information (Bertini et al., 2020). However, in general, successful data visualisations will convey pertinent information in a visually- and cognitively-comprehensible manner (Mackinlay, 1986; Wijk, 2005). Failing to meet these criteria risks misleading viewers, which is antithetical to the purpose of data visualisation. Therefore, knowledge of human factors in visualisation is vital for ensuring charts, graphs, and maps achieve their potential.

## 1.1 Research Motivation and Objective

Interpreting data in any medium does not involve simply observing numerical values, but rather identifying patterns and making *inferences* about the data. For example, one may notice that values are decreasing rapidly, or that there is substantial variability in the dataset, or that some data are missing. One may also make inferences about how large or small values are. This is an important aspect of understanding data, because the same numerical value can be considered large or small depending on its context.

The BBC radio programme *More or Less*, which examines statistics reported in the news and elsewhere, often addresses this issue. Some figures may instinctively seem large, others small. However, asking the question ‘Is it a big number?’ considers whether this initial impression is appropriate. For example, a country’s multi-trillion dollar national debt may appear high, but its level of debt may be similar to that of other countries when accounting for its high GDP (gross domestic product). Similarly, understanding the context of the Richter scale is required to determine that a value less than one does indeed reflect a small earthquake magnitude. Gauging magnitude is important for interpreting numbers in a wide variety of situations.

Because data visualisations are used to convey numerical information, it is important to understand how they may influence judgements of the magnitude of presented values. Studying the cognitive processing of magnitude can reveal how inferences are generated and provide insight into the effects of data visualisation design choices. This, in turn, can inform recommendations for designers who may wish to represent the magnitude of values using graphical cues. Yet, this has been an under-explored topic, with insufficient empirical research exploring this aspect of interpreting visualisations. The objective of this thesis is to generate robust empirical evidence on the interpretation of magnitude in data visualisations.

## 1.2 Contributions

Through a series of empirical experiments, I demonstrate that data visualisation design choices can affect mental representation of numbers’ magnitudes. These large-sample,

controlled experiments strengthen and expand the existing evidence base on this overlooked aspect of data visualisation. I reveal that, in various types of visualisations, graphical cues to context play a role in the processing of how large or small values are. Specifically, judgements about magnitude are informed by the relative positions of values within axis limits. Focusing on underlying cognitive mechanisms generates findings which are applicable to a variety of visualisation formats. These findings also contribute recommendations for designers, which involve considering suitable axis ranges in order to convey magnitude appropriately. This guidance challenges a convention in data visualisation design and advocates against the use of particular default settings, where appropriate.

### **1.3 Overview of the Thesis**

In Chapter 2, I provide context for the empirical work conducted in this project, by reviewing related research. I advocate for an evidence-based approach to data visualisation design which takes into account the mental processing of information. I discuss perceptual and cognitive biases in the interpretation of visualisations, and with a particular focus on the role of numerical axes. This chapter justifies the focus of my investigation and explains how it is informed by prior work in this area.

In Chapter 3, I discuss the methodology and epistemological approach employed in this thesis. In addition to explaining my chosen experimental and statistical techniques, I describe how research projects can benefit from efforts to increase transparency. Approaches for increasing the reproducibility of research are presented, in line with recommendations from a variety of disciplines. A particular emphasis is placed on computational reproducibility: the capacity to recreate the computational environment used in generating results. This provides background for the practices used in the following empirical research chapters.

In Chapter 4, I present a set of three experiments which establish that interpretations of magnitude can be influenced by data visualisation designs. The first experiment demonstrates that manipulating axis limits in dot plots affects participants' judgements of overall magnitude. Two further experiments investigate whether this occurs because axis limits alter the absolute or relative positions of plotted values within axis limits. In dot plots with inverted y-axes, where higher numerical values are presented at lower positions, values near bottom were associated with higher magnitudes. This illustrates that interpretations of magnitude are informed by the relative positions of values within axis limits.

In Chapter 5, I present an experiment which explores how visualisations may influence interpretations of magnitude even when the appearance of plotted values remains unchanged. This experiment demonstrates that manipulating colour legend limits in choropleth maps affects participants' judgements of overall magnitude. Participants rated magnitudes as lower when the range of values on the colour legend extended beyond the largest plotted value. This illustrates that the numerical context accompanying plotted values can influence interpretations of magnitude, without altering the physical appearance of those values.

In Chapter 6, I present a set of two experiments which investigate the role of contextual information on interpretations of magnitude. The first experiment demonstrates that participants' judgements of overall magnitude were affected by extension of bar charts' upper axis limits which incorporated a denominator value. A second experiment reveals that participants' bias was increased when this denominator information was excluded from the text accompanying the chart. This illustrates that knowledge about a dataset's characteristics (e.g., denominator value) can influence the extent to which design choices affect interpretations of magnitude.

Finally, in Chapter 7, I present a synthesis of this empirical work, alongside a discussion of implications and future directions.

## References

- Bertini E, Correll M, Franconeri S. 2020. [Why Shouldn't All Charts Be Scatter Plots? Beyond Precision-Driven Visualizations](#). *arXiv:200811310 [cs]*.
- Franconeri SL, Padilla LM, Shah P, Zacks JM, Hullman J. 2021. The Science of Visual Data Communication: What Works. *Psychological Science in the Public Interest* **22**:110–161. doi:[10.1177/15291006211051956](#)
- Mackinlay J. 1986. Automating the design of graphical presentations of relational information. *ACM Transactions on Graphics* **5**:110–141. doi:[10.1145/22949.22950](#)
- Scaife M, Rogers Y. 1996. External cognition: How do graphical representations work? *International Journal of Human-Computer Studies* **45**:185–213. doi:[10.1006/ijhc.1996.0048](#)
- Wijk JJ van. 2005. The Value of Visualization 8.

## Chapter 2

# Literature Review

### 2.1 A Brief History of Data Visualisation

Throughout history, data visualisations have provided insights on the dominant topics of the day, from science and healthcare to civil rights and warfare. Identifying the first use of data visualisation is impossible, but it is clear that humans have used graphic forms to display numerical information for millennia. For example, on a clay tablet dating from 3100-3000 BC, circles and semicircles represent the quantities of the beer rations which were used to pay workers (MacGregor, 2010). Other early visualisations include geographical maps and astronomical diagrams plotting the movements of the planets. The 18th Century saw the development of many common formats used today, such as bar charts, line charts, and pie charts, all of which are typically credited to William Playfair (Friendly, 2006). However, the late 19th Century has been described as 'The Golden Age of Statistical Graphics' (Friendly, 2008, pg. 13), generating innovations in the representation of large datasets.

In 1855, John Snow produced a map showing the spatial distribution of cholera deaths in an area of London by displaying a mark at the location where each victim had lived (Figure 2.1). Deaths clustered near a contaminated water pump substantiated his radical claim that infected water sources spread this disease (Friendly, 2006). This illustrates how data visualisations can be used to demonstrate vitally important patterns and relationships that were previously overlooked. In 1857, Florence Nightingale visualised fatalities in the Crimean war (Figure 2.2), using a design known as a 'coxcomb', or 'rose diagram' (Friendly, 2006). Each month's death toll was represented by the size of a segment projecting from the chart's centre point. Crucially, the use of colour to distinguish between

different causes of death illustrated that unsanitary conditions in hospitals were a far bigger threat to life than the battlefield (Friendly and Andrews, 2021). This data visualisation was distributed widely to politicians, including the Prime Minister, promoting awareness of the magnitude of preventable deaths (Magnello, 2012). In 1861, Charles Joseph Minard plotted Napoleon's Russian invasion and subsequent retreat with an increasingly diminishing army (Figure 2.3). Part map, part flow diagram, and part line chart, this data visualisation is a paragon of information density, representing six variables in a single graphic whilst telling a coherent story (Tufte, 1986).

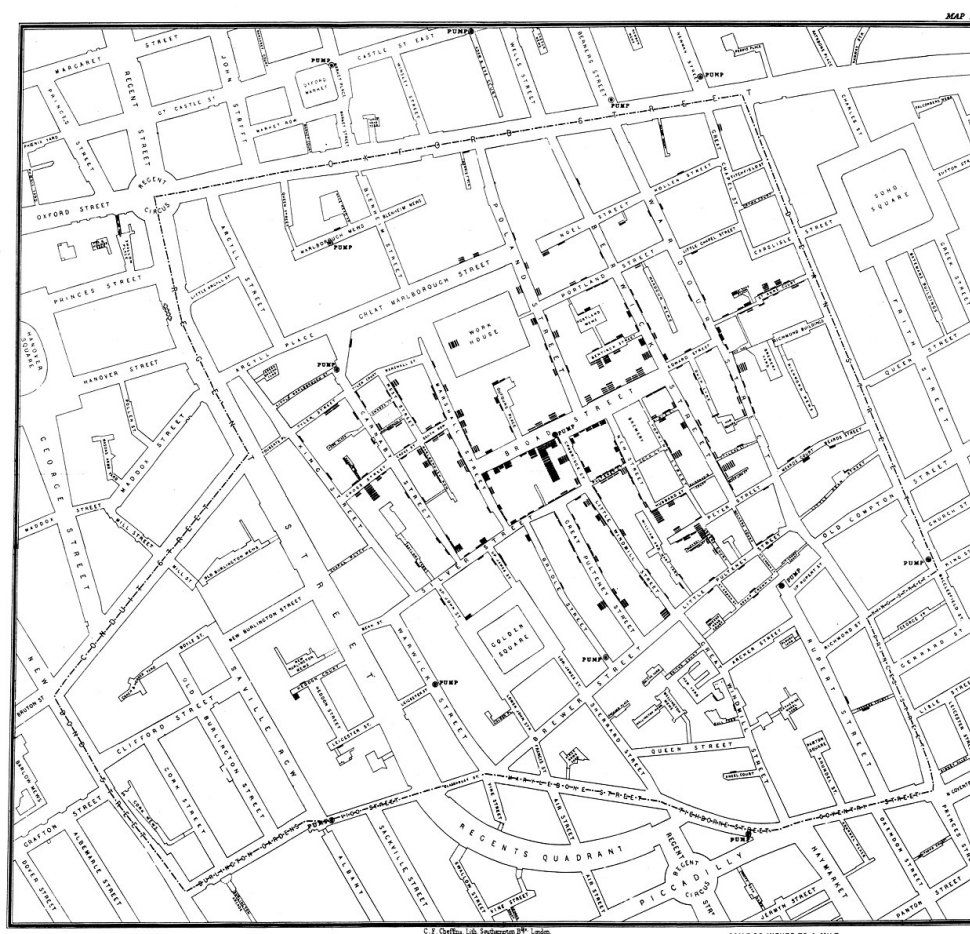


FIGURE 2.1: John Snow's Map of Cholera Deaths (1855)

Although their visualisations may appear to reveal major findings for the first time, none of Snow, Nightingale, or Minard used these visualisations to perform their initial analysis. Instead, these visualisations were used for the purposes of persuasion and storytelling (Kosara and Mackinlay, 2013). This is a testament to the effective use of data visualisations as rhetorical devices and instruments for storytelling, rather than their use as analytical tools. Furthermore, historically significant data visualisations have not always





achieved the recognition and response they sought at the time. W.E.B. Du Bois' data visualisation exhibit on the oppression and development of Black Americans won prizes and medals at the 1900 Paris Exposition (Du Bois, 1900), but was generally overlooked by the mainstream American press (Forrest, 2018).

It is also necessary to acknowledge that the history of data visualisation is rather sparse, and to recognise *contemporary* work in this discipline (Kosara, 2016). Recent innovations in software have generated visualisations with interactive or dynamic elements (Friendly, 2006), but straightforward static visualisations have not disappeared. Indeed, one particularly successful case is the powerfully simple 'warming stripes' visualisation Figure 2.4. This design uses coloured stripes to display average global temperature from 1850 to the present, highlighting the rapid increase in recent years using increasingly darker reds. By eschewing date labels, text, and a colour legend, only the fundamental message remains. Accordingly, this visualisation has been reproduced in various unlikely settings for a data visualisation (e.g., music festivals, clothing, Kintisch, 2019), earning a reputation as a recognisable symbol of the climate emergency.




FIGURE 2.4: Warming Stripes by Ed Hawkins, showing average global temperature data from 1850 to 2018.

When considering famous data visualisations, both historical and contemporary, it is important to avoid making unfounded conclusions about how particular design choices may have contributed to their success. The effectiveness of these designs is undeniable, on account of their documented influence. However, whilst these examples illustrate that visualisations *can* be extremely effective, case studies alone do not provide insight into *why* they are effective. The history of data visualisation reveals the power of visualisations in communication rather than the principles of good design, and speculation about potential positive attributes is not a reliable source of knowledge. This illustrates the importance of studying data visualisations from a scientific perspective.

There is no guarantee that a well-received visualisation unanimously employs effective practices. This is illustrated by Hans Rosling's presentations on global health data (e.g., Hans Rosling, 2006). Research conducted subsequently has revealed variation in the efficacy of his different communication techniques. These hugely popular presentations included verbal explanations of complex animated graphics, delivered with enthusiasm and a dynamic stage presence. Empirical research using Rosling's talks as stimuli demonstrates that his narration facilitates comprehension of data visualisations (Obie et al., 2019). However, the same study found that it has no effect on memory and can elicit concerns about trustworthiness. Another study found that static visualisations of the same dataset improved understanding, but animated visualisations were more popular (Robertson et al., 2008). Furthermore, Rosling's designs also employed variable dot sizes in scatterplots, which can lead to perceptual biases (Anderson et al., 2021; Hong et al., 2021). With so many variables involved in these presentations, more research is required to understand the components of effective storytelling with data visualisation (Kosara and Mackinlay, 2013). Rosling's contributions, in particular his pioneering use of narrative visualisation and his concern for intelligibility, should not be overlooked. Despite this, insight into the *effectiveness* of specific visualisation practices is best acquired through systematic study.

## 2.2 Data Visualisation Formats and the Grammar of Graphics

There is no *one* way to represent a dataset visually. Developing a data visualisation involves making a large number of design choices, which can culminate in vastly different results. Chart 'types' (e.g., bar chart, pie chart, line chart) offer an easy way to categorise the broad format of a visualisation. However, these categorisations do not reflect the way that data visualisations are constructed or how they function. The 'Grammar of Graphics' (Wilkinson, 2005) offers an alternative approach.

The Grammar of Graphics is a system for formally defining visualisations in terms of their underlying structure. As a *grammar*, rather than a taxonomy, it was developed in order to express the composition of any data visualisation through six components (Wilkinson, 2005). *Elements* describe both the *aesthetic attributes* which visually encode values (e.g., position, size, hue, transparency), and the *geometries* which represent those values (e.g., bar, dot, line). *Coordinate systems* describe the canvas used for representing values. For example, Cartesian coordinates use the vertical and horizontal dimensions associated

with bar charts; polar coordinates use the circular mapping associated with pie charts; and map projections use a cartographic mapping associated with world maps. The other components of the Grammar of Graphics are the *data* used, *variable transformations* (e.g., mean, sum, rank), *scale transformations* (e.g., linear scale, logarithmic scale) and guides (e.g., axes, colour legends).

This system allows for efficient and consistent characterisation of different visualisation formats. For example, Figure 2.5 shows that bar charts (A) and dot plots (B) both use the same aesthetic attribute (position) to encode values, but differ in their geometry (bar versus dot). A regular bar chart (C) is equivalent to a chart like Florence Nightingale's coxcomb (D), when using polar coordinates. Conversely, a *stacked* bar chart (E) is equivalent to a regular pie chart (F), when using polar coordinates. Visualisations can employ more than one aesthetic attribute, for example, the examples in Figure 2.5 use hue to represent the different categorical values. However, it is possible to use lightness to represent different numerical values instead, with darker colours representing higher values (G). Using this aesthetic attribute, in combination with a map projection and geometries based on the shape of geographic regions, produces a choropleth map (H). This illustrates how the components can be combined in a flexible and modular manner, with many more possible visualisations of this dataset. The Grammar of Graphics has been influential in the development of a number of data visualisation design tools, including Polaris (Stolte et al., 2002), which became Tableau, ggplot2 in R (Wickham, 2010), D3 (Bostock et al., 2011), and Vega-Lite (Satyanarayan et al., 2017).

## 2.3 Data Visualisation Software and the Influence of Default Settings

Considering the *process* by which data visualisations are created is crucial for understanding this subject. Modern software has made it possible to quickly and easily produce a wide range of visualisations. However, variation across visualisation design tools affects the range of visualisation formats available to users and the degree of customisation offered. For example, programming libraries, where data visualisations are created by writing lines of code (e.g., ggplot2 in R, plotly in Python) typically offer more options and greater control than simple point-and-click software (e.g., Microsoft Excel). Different visualisation design tools also provide different specialised capabilities. For example, Tableau is often

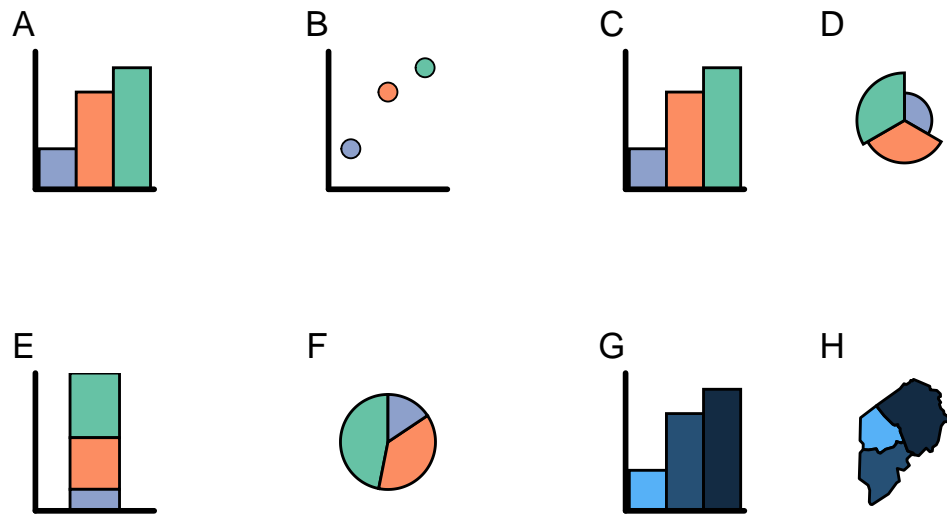


FIGURE 2.5: Eight different ways of displaying the same dataset, to illustrate the Grammar of Graphics.

used for business intelligence applications, such as building dashboards (Elias et al., 2013), D3 was developed for designing visualisations for the web (Bostock et al., 2011), VegaLite was developed for generating interactive visualisations (Satyanarayan et al., 2017), and ggplot2 was developed for use within a data analysis workflow (Wickham, 2011).

What is possible and practicable using a particular piece of software will influence a designer's choices, which may in turn affect viewers' interpretations. However, visualisation software also initially imposes particular properties on a visualisation, through its default settings. For example, a pie chart, prior to customisation, will present segments in a particular order, using a particular set of colours. Even when it is *possible* to reject default designs, they can be highly influential, because they will remain unchanged when it is unclear *how* or *why* they should be altered (Shah and Kesan, 2006).

However, existing default settings in data visualisation are not always suitable. For example, by default, software for creating line charts typically employs y-axes which are constrained to the range of the data. Therefore, the highest and lowest values are presented at the chart's extremes, impeding viewers' ability to gauge the magnitude of the difference. When representing categorical values with colour, designers can improve viewers' performance by rejecting defaults in favour of colours which correspond semantically

with plotted data [e.g., yellow for banana; Lin et al. (2013)]. When representing continuous values with colour, rainbow colour palettes are a popular choice (Ware et al., 2023). However, alternative colour palettes have been employed as defaults in attempts to avoid issues in perception associated with the rainbow colour palette (Reda and Szafir, 2021).

Default settings can also cause issues in the generation of visualisations for exploratory analysis (Correll et al., 2019). These visualisations are used to understand the characteristics of a dataset prior to formal statistical analysis. For example, histograms are used to visualise the shape of a univariate distribution. The algorithm used to produce histograms in R and D3 assumes by default that the data are normally-distributed. Consequently, abnormalities in non-normal data can be ‘smoothed-over’, preventing viewers from identifying them. Dot plots, another format for visualising distributions, display each individual value using a dot. In R and Tableau, by default, these dots have no translucency. This can result in many overlapping dots in close proximity, impeding a viewer’s ability to differentiate between areas with different densities. This demonstrates that default settings are not *always* inappropriate. However, when default settings are *agnostic* regarding the characteristics of plotted data, unquestioning use of them can conceal relevant aspects of a dataset.

Although the above research exposes issues with some default settings, they are certainly not exclusively harmful. For example, one default setting used by Microsoft Excel is ‘redundant encoding’, where individual data points are represented using different shapes *and* different colours (e.g., blue diamonds and green triangles). Experimental work has observed that whilst this technique does not confer benefits in some tasks (Gleicher et al., 2013), it improves viewers’ performance in other tasks (Nothelfer et al., 2017). Empirical research is important in order to identify how default settings may be beneficial or detrimental. Indeed, researchers in data visualisation often suggest that their findings may inform the development of default settings (e.g., Heer and Bostock, 2010; Kerns and Wilmer, 2021; Xiong et al., 2021). Several experiments in this thesis were designed to explore the consequences of default settings.

## 2.4 Popular Guidance on Effective Data Visualisation Design

Our understanding of how people interpret data visualisations (and subsequent guidance) is built on shaky foundations. Some received wisdom has not been empirically tested

at all, other claims have been discredited or confirmed only recently (Kosara, 2016). Consequently, it is not always clear where evidence ends and opinion starts; intuition and unsubstantiated statements make for “visualisation folklore” (Correll, 2022, pg. 3).

Statistician Edward Tufte is a source of widely-cited advice on the design of data visualisations, which he has articulated in popular books such as *The Visual Display of Quantitative Information* (Tufte, 1986). One famous contribution is the ‘lie-factor’, which attempts to quantify the degree of misrepresentation in charts that distort data. For example, plotting values using two dimensional images exaggerates differences between values, because perceived size is determined by an image’s entire *area*, not just its *height*. Consequently, in Figure 2.6, which appears to show a 42 percentage point difference, dividing by the real numerical difference of 15 percentage points generates a lie-factor of 2.8, compared to an ideal score of 1. However, Tufte’s criteria proposed for diagnosing *substantial* distortion (less than 0.95 or more than 1.05) are based on speculation, rather than scientific evidence (Beattie and Jones, 2002). In similarly arbitrary guidance, Tufte (1986) suggests that a dataset of 20 or fewer observations should be presented in a table, rather than a data visualisation. However, a subsequent empirical experiment has revealed that pie charts elicited more accurate responses than tables for proportion judgements involving only three observations (Spence and Lewandowsky, 1991).

Tufte also advocates for minimalism in the design of data visualisations. His recommendation to maximise the ‘data-ink ratio’ involves maximising the proportion of ink (i.e., pixels) used to depict the data itself and minimising inessential elements (Tufte, 1986). However, this notion is vague and prone to excessive simplicity. Redundant features can serve to minimise error (Tversky, 1997), with ‘redundant’ tick marks on axes required for accurately extracting numerical values (Kosslyn, 1985). The qualifier “within reason” (Tufte, 1986, pg. 96) is an imprecise addition to this guidance, whereas empirical research can identify where extreme sparseness unnecessarily biases interpretations (Gillan and Richman, 1994; Stock and Behrens, 1991).

Consistent with his minimalistic approach, Tufte’s recommendation to eliminate ‘chartjunk’ involves avoiding the use of distracting visual embellishments, which range from excessive gridlines to artistic decoration (Tufte, 1986). However, there is mixed evidence regarding the harm caused by ‘chartjunk’ (Franconeri et al., 2021). However, condemning chartjunk remains popular, not just on aesthetic grounds, but also due to




FIGURE 2.6: A reproduction of a data visualisation used by Tufte (1986) to illustrate the 'lie-factor', originally published in the Los Angeles Times. The final value of 12% is a 15 percentage point decrease from the first value of 27%. However, Tufte's lie-factor of 2.8 indicates that this reduction is perceived as a 42 percentage point decrease.

the rhetorical qualities of minimalist designs, which imply a straightforward, unbiased presentation of data (Kennedy et al., 2016; Kosara, 2016).

Researchers argue that Tufte's recommendations for minimalistic designs do not account for human cognitive processing (Chabris and Kosslyn, 2005; Wilkinson, 2005). More generally, he has been criticised for failing to support his claims with empirical evidence (Feldman-Stewart et al., 2000). Instead, his guidance is underpinned by a large collection of example visualisations taken from various sources. Therefore, Tufte's principles might assist in describing common features of some successful visualisations, rather than serving as definitive rules (Kindlmann and Scheidegger, 2014). Rigorous data visualisation research is required to fill gaps in knowledge and generate a reliable evidence-base.

## 2.5 Empirical Research on Data Visualisation

Visualisation research takes many forms. Empirical studies on data visualisation have employed a range of techniques, including controlled experiments, usability tests, interviews, observations, and case studies, and have focused variously on perception, cognition, exploratory data analysis, and user experience (Lam et al., 2012). Experimental psychology studies on data visualisation are particularly valuable because they generate fundamental evidence on *how* visualisations are interpreted. Considering human interpretation in visualisation research is crucial for generating generalisable knowledge. Inadequate best practice recommendations indicate insufficient understanding of psychological mechanisms. However, progress can be slow, since theories about cognitive and perceptual processes are built through cumulative work (Chen et al., 2020). Existing psychological research confers benefits in the form of related empirical work, as well as contributing established methods and theories (Correll, 2022; Rensink, 2021).

Several studies illustrate that preferences and introspection are not always a reliable source of information on effective visualisation practices. For example, an experiment exploring physicians' judgements about clinical trials found that icon arrays resulted in the most accurate judgements, compared to tables, pie charts, and bar charts (Elting et al., 1999). However, none of the 34 physicians in the sample preferred this format. In another study, medical students almost unanimously preferred visualisations with a rainbow colour scheme, but made fewer errors when using a diverging (e.g., red-blue)



colour scheme (Borkin et al., 2011). Furthermore, tables of values are favoured over visualisations in certain tasks where visualisations actually offer significant benefits (Saket et al., 2019) and certain statistical map designs are preferred over others despite conferring no performance advantages (Mendonça and Delazari, 2014). Participants in Burns et al.'s (2021) study estimated that pictographs took longer to understand compared to equivalent visualisations without icons. However, this self-report measure was at odds with recorded response times, which indicated no differences between visualisations types. Many authors suggest that preferences are influenced by familiarity, rather than effectiveness. Measuring preferences provides valuable insight into people's engagement with different visualisations. However, such opinions must be treated appropriately, not used to inform conclusions about efficacy.

Rensink (2021) presents recommendations for generating useful findings in data visualisation research. Using a single task, and manipulating a single feature of interest, over multiple trials, assists in identifying underlying mechanisms. Integrating explanations from prior research helps ensure explanations of mental processes are sufficiently detailed. Other important but frequently overlooked matters include appropriate counterbalancing, reporting effect sizes, and acknowledging individual differences.

There are a multitude of variables that can be manipulated to gain insight into visualisations. Criticisms are sometimes levelled at studies with particularly high or low levels of experimental control. However, researchers must strike an appropriate balance between ecological validity and precision (Abdul-Rahman et al., 2020). Choosing suitable tasks for participants requires a similar trade-off (Suh et al., 2022).

Vision sciences offer a variety of paradigms for assessing various aspects of human performance in visualisation tasks. For example, experiments may evaluate accuracy (by comparing responses to a correct answer), precision (by quantifying variability in responses), or processing speed (by measuring reaction times, Elliott et al., 2020). However, chosen methods must be appropriate for a research question. Whereas methods from vision-sciences are typically concerned with performance in low-level perceptual tasks, other research focuses on decision-making (Padilla et al., 2018) or *message*-level interpretations (Pandey et al., 2015). The latter concerns overall assessments of data, such as whether a difference is large or small, rather than the ability to extract specific values. This is also referred to as *gist* (Reyna and Brainerd, 1991).

## 2.6 Perceptual Precision in Data Visualisations

Identifying gaps in our understanding of the psychology of data visualisations requires knowledge of prior lines of inquiry and established findings. Arguably the most influential study in the field of data visualisation is Cleveland and McGill's (1984) investigation of elementary perceptual processes involved in viewing visualisations. This study sought to establish how *precisely* viewers can represent different graphical properties used to encode data (e.g., position, length, angle, etc.). For each encoding type, participants identified which of two marks conveyed the smaller value, and estimated the difference in size as a percentage. Subsequent ranking based on the magnitude of participants' errors produced a hierarchy of visual encoding channels. Since position-encoding produced smaller errors than both length- and angle-encoding, this suggests that viewers will represent data most precisely when it is encoded using position on a common (aligned) scale.

This study's findings have endured replication (Heer and Bostock, 2010) and enthusiasm for perceptual precision has inspired a great deal of important research in this field. This research spans visual processing of proportion (Hollands and Spence, 1998; Spence and Lewandowsky, 1991), variance (Stock and Behrens, 1991), correlation (Harrison et al., 2014; Hong et al., 2021), and other basic processes, such as visual comparison (Simkin and Hastie, 1987; Zacks et al., 1998) and colour discrimination (Szafir, 2018). The study has also influenced development of software for automating visualisation design (Mackinlay, 1986) and simulating visualisation comprehension (Lohse, 1993). However, to consider perceptual precision as the *only* relevant concern in data visualisation design is unwarranted; many additional factors require consideration.

## 2.7 Beyond Perceptual Precision

Optimally-precise visual cues are not always employed when viewing visualisations. Viewers are sensitive to other task-irrelevant visual cues, which can lead to inaccurate judgements about plotted data (Yuan et al., 2019). In particular tasks, precision can actually hinder, rather than facilitate, judgements. For example, because perceptual averaging benefits from a lower spatial frequency, less precise colour encoding offers *greater* efficiency than more precise position encoding (Correll et al., 2012), see Figure 2.7. Furthermore, effective decision-making under uncertainty does not necessarily correspond to precision

in probability estimation, because of the differences in mental processing associated with these two distinct tasks (Kale et al., 2020).




FIGURE 2.7: Visualisations used to demonstrate greater perceptual averaging ability for colour encoding (low spatial frequency), than position encoding (high spatial frequency). Participants more easily identified the month with the highest average temperature (August), in the colour bar visualisation, than a line chart showing the same data. Adapted from Correll et al. (2012).

Furthermore, the choice of graphical encodings employed in a data visualisation can influence the *type* of interpretation it elicits. For example, viewers are more likely to refer to trends when describing line graphs and discrete differences when describing bar charts (Zacks and Tversky, 1999). Similarly, *production* of bar charts and line charts is also influenced by whether a discrete or continuous relationship is specified in the brief. Design choices also influence beliefs about the distribution of underlying data, when presenting average values (Newman and Scholl, 2012). Compared to a data point positioned ‘outside’ a bar, a data point positioned ‘inside’ a bar is more likely to be considered part of the underlying data (see Figure 2.8). However, displaying only confidence intervals eliminates this bias (Pentoney and Berger, 2016). This accords with the notion that viewers’ cognitive associations between visual features and abstract characteristics of data are important in data visualisation design. Through common metaphors (e.g., hierarchy and vertical position), aspects of a design may offer *affordances*, carrying connotations which encourage particular interpretations (Kindlmann and Scheidegger, 2014; Xiong et al., 2022; Ziemkiewicz and Kosara, 2008).

Attention is another important factor in the comprehension of data visualisations. Complex tasks requiring selective attention can cause distinctive patterns in non-focal data to be completely overlooked (Boger et al., 2021). Features of data mentioned in textual summaries are over-weighted in viewers’ mental representations, causing difficulty in the




FIGURE 2.8: Examples of data points positioned ‘inside’ and ‘outside’ a bar showing an average value, and the equivalent plots with confidence intervals. The data point (grey circle) more likely to be considered part of the underlying data in A than B, but is equally likely to be considered part of underlying data in C and D. Adapted from Newman and Scholl (2012) and Pentoney and Berger (2016), from simplified versions of experimental stimuli.

ability to assume the perspective of a naïve viewer (Xiong et al., 2019). In addition, the salience of vertical bars may be responsible for erroneous judgements of differences between histograms with identical distributions (Lem et al., 2014). As a solution, explicitly encoding differences between pairs of values can facilitate recognition of relevant patterns (Nothelfer and Franconeri, 2020) and highlighting particular attributes can facilitate recall (Ajani et al., 2021).

Simply conveying information is not the only purpose of data visualisations, since they also influence recall, opinion-formation, and decision-making (Bertini et al., 2020). As illustrated above, a large number of cognitive biases affect these aspects of the mental processing of data, as well as several others, including causal reasoning and assessment of hypotheses (Dimara et al., 2020). Whilst it is necessary to consider the precision of elementary perceptual processes, that alone is not sufficient for a comprehensive understanding of how data visualisations function (Bertini et al., 2020).

## 2.8 Manipulating Axes in Data Visualisations

Understanding how inaccurate impressions arise provides insight into the mechanisms involved in interpreting data visualisations. This, in turn, can inform recommendations for effective design. A prominent topic in the literature on misleading visualisations is axis truncation. This typically refers to the practice of employing a y-axis which commences with a non-zero value (Correll et al., 2020) though may also be considered any reduction at either extreme of an x- or y-axis (Pandey et al., 2015). Figure 2.9 shows examples of truncated and non-truncated y-axes in line and bar charts.




FIGURE 2.9: Examples of truncated and non-truncated line and bar charts, showing the same data. The visual difference between values appears larger in the truncated version, compared to the non-truncated version due to the smaller range of values on the y-axis.

There is considerable evidence that axis limits influence interpretations of data. The majority of research on this topic has focused on how constraining the range of an axis, and thus increasing the physical distance between plotted values, increases the perceived magnitude of the difference between those values. For example, one study reports that accountants appraising financial performance using line and bar charts interpreted plotted increases as larger when these increases were depicted using a truncated y-axis (Taylor and Anderson, 1986). Similarly, bar charts employing truncated axes biased students' investment decisions (Arunachalam et al., 2002). Students were more likely to select a less-successful company when a truncated chart exaggerated that company's growth rate,

compared to when a non-truncated chart was used. An online experiment also observed that differences between values were considered larger when truncated bar charts were used (Pandey et al., 2015). This experiment examined message-level representations of data by framing questions in terms of subject matter (e.g., access to safe drinking water) rather than graphical elements (e.g., difference in bar length). Other axis manipulations, such as log-scales (Romano et al., 2020), inverted scales Pandey et al. (2015), and expanded axes in scatterplots (Cleveland et al., 1982) also influence judgements about data.

Risk communication research has independently generated similar findings about axis truncation. Because many hazards cannot be completely avoided, data visualisations are often used to contrast the levels of risk associated with two scenarios (e.g., an intervention vs. no intervention). Thus, assessments of 'risk reduction' are essentially judgements about the magnitude of difference between two values. For example, one experiment compared stacked bar charts, which include additional information on the total number of individuals at risk, to bar charts displaying only the number of individuals *affected* (Stone et al., 2003). The latter design increased the bars' visual disparity, and subsequently increased impressions of the magnitude of difference.

The physical distance between data points consistently biases interpretations of the magnitude of difference in spite of attention to actual numerical values and also design features intended to highlight truncation (Correll et al., 2020). Bias is diminished, but still observed, following explicit warnings about errors in judgement due to y-axis truncation. This suggests that this effect is largely automatic, and does not primarily occur due to insufficient engagement of cognitive capabilities (Yang et al., 2021).

Researchers have also explored individual differences in interpretations of data presented using truncated axes. One study observed no association between participants' susceptibility to bias due to axis truncation in bar charts and their data visualisation literacy (Yang et al., 2021). Conversely, another experiment suggests that the effect of axis truncation on subjective judgements and quantitative estimates in line charts disappears when accounting for data visualisation literacy (Driessen et al., 2022). However, in the latter experiment, low variability in observed data visualisation literacy levels raised concerns about the sensitivity of the scale used to measure data visualisation literacy.

Pandey et al. (2015) and Yang et al. (2021) propose that this bias could arise due to

the dominance of first impressions during translation from graphical schemata (Pinker, 1990) to a ‘real-world’ conceptual understanding Tversky and Kahneman (1974). Additionally, Yang et al. (2021) suggest that viewers’ beliefs about the communicative intent of a designer could play a role in viewers’ interpretations. Under Grice’s *Co-operative Principle* (Grice, 1975), communicative contributions in conversation are assumed to be truthful, relevant, clear, and sufficiently informative. Extrapolating this principle to data visualisations, viewers might infer that differences between values must be genuinely large if they appear large, because they would otherwise not be presented as such.

In *How to Lie With Statistics*, Huff (1993) suggests that axis truncation creates a false impression of plotted data. This practice has been labelled ‘deceptive’ for both bar and line charts (Lauer and O’Brien, 2020). Furthermore, a tool for automatically identifying and correcting misleading line charts extends y-axes to include zero whenever this value is omitted from the original chart (Fan et al., 2022).

Recent work has presented an alternative perspective on this controversial practice. Non-truncated axes can obscure significant differences just as easily as truncated axes can exaggerate inconsequential differences. The appropriate magnitude to convey depends on what constitutes an important difference in the data at hand (Correll et al., 2020). Indeed, *failing* to truncate an axis could be considered misleading in certain circumstances (Wainer, 1984). Yang et al. (2021) suggest that effective designs will ensure that a viewer’s immediate characterisation of plotted data closely corresponds to their interpretation following a detailed inspection. Acknowledging that differences must be depicted in proportion to their significance, (Witt, 2019) reports that axes spanning approximately 1.5 standard deviations provide a balance between sensitivity and bias in fields with standardised effect size measures, such as psychology. Unfortunately, different domains will not necessarily share the same notion of what amounts to a meaningful difference. Choices regarding axis ranges are ultimately designers’ unavoidable decisions (Correll et al., 2020).

Finally, although line charts and bar charts are equally susceptible to biases due to truncation (Correll et al., 2020; Witt, 2019), there may be reason to treat them differently. Truncation distorts the mapping between a bar’s extent and the quantity it represents, but the free-floating position-encoding used in line charts does not convey quantity in the same manner, providing immunity against such distortion (Bergstrom and West, 2017). Therefore, whilst starting a bar chart’s axis at zero cannot guarantee that differences

between values are depicted appropriately, this does ensure adherence to a fundamental aspect of visualisation design. Alternatively, to avoid this trade-off, quantitative data with discrete categories can be plotted using position-encodings only (e.g., dot plots rather than bar charts).

## 2.9 Misleading Data Visualisations

Some misleading visualisations prevent viewers from accurately extracting numerical information. However, research on axis truncation illustrates that misleading visualisations may also interfere with *subjective* judgements. A line chart may precisely represent a dataset's numerical properties yet generate a distorted impression of the magnitude of a trend. The latter is revealed not by assessing viewers' *performance*, but their *interpretations* (Stone et al., 2015).

Influencing subjective judgements may still be considered a *misleading* practice because a dishonest framing of information could elicit an unreliable interpretation which would differ from the same viewer's better-informed perspective. For example, an increase may initially appear small, but may be interpreted as large when depicted in the context of genuinely meaningful differences. Not all aspects of deceptive design are *inherently* misleading, and deceptiveness can be context-dependent. Comparing examples of 'misleaders' from Ge et al.'s (2023) design space helps illustrate this distinction. 'Concealed uncertainty' and 'cherry-picking' refer to unambiguously deceptive practices, whereas 'aggregation' and 'scale range' must be preceded by the word *inappropriate* in order to convey their capacity to deceive.

## 2.10 Data Visualisation Literacy

Understanding individual differences in the ability to comprehend data in visualisations is important for understanding the psychology of data visualisations (Boy et al., 2014). Research on this topic requires reliable tools for measuring data visualisation literacy.

Galesic and Garcia-Retamero's 13-item test (2011) is based on Friel et al.'s (2001) hierarchy of skills for interpreting visualisations, which ranges from comprehension to extrapolation. Research has demonstrated that this scale helps to predict whether a graphical



representation will facilitate understanding of risk information (Okan et al., 2012a). A different 53-item test employs a wide range of data visualisation formats, and higher scores are positively associated with both numeracy and need for cognition (Lee et al., 2019).

Research on data visualisation literacy has tended to focus on interpretation of well-designed charts (Ge et al., 2023). However, the ability to detect (Camba et al., 2022) and make sense of (Ge et al., 2023) misleading charts should be considered an important feature of data visualisation literacy. A robust 30-item test enables assessment of an individual's ability to accurately comprehend deceptive designs (Ge et al., 2023). This work also suggests that attention and critical thinking may benefit viewers in avoiding some, but not all, biased interpretations. Furthermore, using Galesic and Garcia-Retamero's 13-item test (2011), Okan et al. (2016) found that higher data visualisation literacy is associated with more time processing a visualisation's misleading features, thus promoting correct interpretations. Lower data visualisation literacy is associated with greater reliance on conventions (e.g., the relationship between vertical position and magnitude).

The empirical work presented in this thesis employs the 5-item version of Garcia-Retamero et al.'s (2016) Subjective Graph Literacy scale. Users are asked to rate their competence in working with bar charts, line charts, and pie charts, and also their ability to perform simple tasks using bar charts. The subjective approach echoes prior work in the development of subjective numeracy scales. Despite its short completion time and use of subjective ratings, it is strongly correlated with an *objective* measure of data visualisation literacy (2011). The scale also produces a final score out of 30, offering greater sensitivity than a similarly brief objective scale, where tallying correct responses produces a final score out of 4 (Okan et al., 2019). These characteristics make for an appropriate tool for assessing participants' data visualisation literacy in experimental studies. Indeed, this measure has been used to assess variability between participants in studies on axis truncation (Yang et al., 2021), correlation (Strain et al., 2023), information synthesis (Mantri et al., 2022), and explanation of visualisations (Yang et al., 2023).

## 2.11 Interpreting Absolute Magnitude

Data visualisation design has the potential to impact subjective judgements about many aspects of data, such as variability, noise, and numerosity. Prior research has closely

examined how axis truncation can influence judgements of *relative* magnitude (i.e., differences between values). In contrast, little is known about how axis limits may influence judgements of *absolute* magnitude: how large or small values are. Despite this, these judgements can be fundamental in developing a basic interpretation of quantitative data. For example, assessing the probability of rain, the number of patients on a waiting list, the amount of CO<sub>2</sub> emitted during a journey, or the level of support for a political candidate, are all judgements of *absolute* magnitude. Limited insight into how magnitude is interpreted in data visualisations impedes understanding of how visualisations may effectively communicate magnitude. Prior research on this topic is summarised below.

In bar charts displaying data on individuals affected by a risk (e.g., Figure 2.10), perceived likelihood decreases when the total population at risk is emphasised using shaded bars (top), rather than blank space (bottom, Stone et al., 2017). In bar charts violating the convention of mapping higher values to higher positions, participants frequently misinterpreted absolute magnitudes (Okan et al., 2012b). This was due to difficulty in rejecting first impressions, particularly for participants with low data visualisation literacy. However, other work which has combined judgements of values' absolute magnitudes with judgements of relative differences has impeded analysis of the former (Okan et al., 2018). Researchers have suggested that visualisations which facilitate comprehension of relative differences may fail to effectively communicate the absolute magnitudes of the values depicted, illustrating a potential trade-off in design (Reyna, 2008).

One study has specifically focused on how *axis ranges* may inform impressions of absolute magnitude. Sandman et al. (1994) manipulated the design of risk ladders, where individual probabilities are presented on vertical scales incorporating a range of probability values. Changing this range alters the position of a plotted value. This study found that perceived threat (a composite measure made up of perceived likelihood, danger, reported concern and fear) was higher when the risk appeared near to the top of the ladder, compared to near the bottom. This finding resembles framing effects described in the psychology literature, wherein interpretations of the same information differ according to the manner in which it is presented (Tversky and Kahneman, 1981). However, the position of plotted values did not completely dictate magnitude judgements. A numerically higher risk plotted at *the same position* near the top of the ladder generated higher ratings. There was also mixed evidence regarding the effects of the axis manipulation on intentions to spend money mitigating the risk. Confidence in the robustness of these findings is limited by




FIGURE 2.10: Stimuli used to demonstrate that including additional shaded bars (grey) decreases interpretations of the magnitude of other values (blue), compared to blank space. Adapted from Stone et al. (2017).

various factors including use of a single trial per participant, a single scenario, a composite measure obscuring pure magnitude ratings, and a confounding variable of the risk ladder's range.

Comparing linear and logarithmic risk ladders, Freeman et al. (2021) did not replicate Sandman et al.'s (1994) main finding. However, in addition to a graphical cue to magnitude, they used risk ladders which employed additional symbolic number cues in their titles, labels, and accompanying descriptions (e.g., "12%", "120 out of 1000"). A broken scale may also have reduced the degree to which inferences were based on the value's physical position. Therefore, participants' judgements may not have been based purely on the appearance of visualisations.

Compared to knowledge on the interpretation of *relative* magnitude in data visualisations, knowledge on the interpretation of *absolute* magnitude is limited. Insufficient research on this topic impedes understanding of the relationship between design choices and subsequent impressions. The cognitive processing of absolute magnitude in data visualisations is the focus of this thesis.

## 2.12 The Present Thesis

This thesis consists of a detailed investigation into how interpretations of the magnitude of numerical values are influenced by design choices. Three sets of experiments (six experiments in total) each address the overall objective of the thesis: to understand the cognitive processing of absolute magnitude in data visualisations. However, each set of experiments uses a different data visualisation format and investigates a different factor influencing judgements, thus contributing to a comprehensive account of cognitive processing. The following chapter introduces the methodology and epistemological approach employed in this work, prior to the empirical experiments in Chapters 4-6.

## References

- Abdul-Rahman A, Chen M, Laidlaw DH. 2020. A Survey of Variables Used in Empirical Studies for Visualization In: Chen M, Hauser H, Rheingans P, Scheuermann G, editors. *Foundations of Data Visualization*. Cham: Springer International Publishing. pp. 161–179. doi:[10.1007/978-3-030-34444-3\\_7](https://doi.org/10.1007/978-3-030-34444-3_7)
- Ajani K, Lee E, Xiong C, Nussbaumer Knaflitz C, Kemper W, Franconeri S. 2021. Declutter and Focus: Empirically Evaluating Design Guidelines for Effective Data Communication. *IEEE Transactions on Visualization and Computer Graphics* 1–1. doi:[10.1109/TVCG.2021.3068337](https://doi.org/10.1109/TVCG.2021.3068337)
- Anderson ER, Maeda K, Yousif SR, Franconeri SL. 2021. Inconsistent Perception for Cumulative Area in Bubble Plots 5.
- Arunachalam V, Pei BKW, Steinbart PJ. 2002. Impression Management with Graphs: Effects on Choices. *Journal of Information Systems* **16**:183–202. doi:[10.2308/jis.2002.16.2.183](https://doi.org/10.2308/jis.2002.16.2.183)
- Beattie V, Jones MJ. 2002. Measurement distortion of graphs in corporate reports: An experimental study. *Accounting, Auditing & Accountability Journal* **15**:546–564. doi:[10.1108/09513570210440595](https://doi.org/10.1108/09513570210440595)
- Bergstrom C, West J. 2017. [Misleading axes on graphs](#).
- Bertini E, Correll M, Franconeri S. 2020. [Why Shouldn't All Charts Be Scatter Plots? Beyond Precision-Driven Visualizations](#). *arXiv:200811310 [cs]*.
- Boger T, Most SB, Franconeri SL. 2021. Jurassic Mark: Inattentional Blindness for a Datasaurus Reveals that Visualizations are Explored, not Seen 5.
- Borkin M, Gajos K, Peters A, Mitsouras D, Melchionna S, Rybicki F, Feldman C,

- Pfister H. 2011. Evaluation of Artery Visualizations for Heart Disease Diagnosis. *IEEE Transactions on Visualization and Computer Graphics* **17**:2479–2488. doi:[10.1109/TVCG.2011.192](https://doi.org/10.1109/TVCG.2011.192)
- Bostock M, Ogievetsky V, Heer J. 2011. D<sup>3</sup> Data-Driven Documents. *IEEE Transactions on Visualization and Computer Graphics* **17**:2301–2309. doi:[10.1109/TVCG.2011.185](https://doi.org/10.1109/TVCG.2011.185)
- Boy J, Rensink RA, Bertini E, Fekete J-D. 2014. A Principled Way of Assessing Visualization Literacy. *IEEE Transactions on Visualization and Computer Graphics* **20**:1963–1972. doi:[10.1109/TVCG.2014.2346984](https://doi.org/10.1109/TVCG.2014.2346984)
- Burns A, Xiong C, Franconeri S, Cairo A, Mahyar N. 2021. Designing with Pictographs: Envision Topics without Sacrificing Understanding. *IEEE Transactions on Visualization and Computer Graphics* 1–1. doi:[10.1109/TVCG.2021.3092680](https://doi.org/10.1109/TVCG.2021.3092680)
- Camba JD, Company P, Byrd V. 2022. Identifying Deception as a Critical Component of Visualization Literacy. *IEEE Computer Graphics and Applications* **42**:116–122. doi:[10.1109/MCG.2021.3132004](https://doi.org/10.1109/MCG.2021.3132004)
- Carpenter PA, Shah P. 1998. A Model of the Perceptual and Conceptual Processes in Graph Comprehension 26.
- Chabris CF, Kosslyn SM. 2005. Representational Correspondence as a Basic Principle of Diagram Design 22.
- Chen M, Abdul-Rahman A, Laidlaw DH. 2020. [The Huge Variable Space in Empirical Studies for Visualization – A Challenge as well as an opportunity for Visualization Psychology.](#)
- Cleveland WS, Diaconis P, McGill R. 1982. Variables on Scatterplots Look More Highly Correlated When the Scales Are Increased. *Science* **216**:1138–1141. doi:[10.1126/science.216.4550.1138](https://doi.org/10.1126/science.216.4550.1138)
- Cleveland WS, McGill R. 1984. Graphical Perception: Theory, Experimentation, and Application to the Development of Graphical Methods. *Journal of the American Statistical Association* **79**:531–554. doi:[10.2307/2288400](https://doi.org/10.2307/2288400)
- Correll M. 2022. [Are We Making Progress In Visualization Research?](#)
- Correll M, Albers D, Franconeri S, Gleicher M. 2012. Comparing averages in time series dataProceedings of the 2012 ACM Annual Conference on Human Factors in Computing Systems - CHI '12. Austin, Texas, USA: ACM Press. p. 1095. doi:[10.1145/2207676.2208556](https://doi.org/10.1145/2207676.2208556)

- Correll M, Bertini E, Franconeri S. 2020. Truncating the Y-Axis: Threat or Menace? Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems. Honolulu HI USA: ACM. pp. 1–12. doi:[10.1145/3313831.3376222](https://doi.org/10.1145/3313831.3376222)
- Correll M, Li M, Kindlmann G, Scheidegger C. 2019. Looks Good To Me: Visualizations As Sanity Checks. *IEEE Transactions on Visualization and Computer Graphics* **25**:830–839. doi:[10.1109/TVCG.2018.2864907](https://doi.org/10.1109/TVCG.2018.2864907)
- Dimara E, Franconeri S, Plaisant C, Bezerianos A, Dragicevic P. 2020. A Task-Based Taxonomy of Cognitive Biases for Information Visualization. *IEEE Transactions on Visualization and Computer Graphics* **26**:1413–1432. doi:[10.1109/TVCG.2018.2872577](https://doi.org/10.1109/TVCG.2018.2872577)
- Driessen JEP, Vos DAC, Smeets I, Albers CJ. 2022. Misleading graphs in context: Less misleading than expected. *PLOS ONE* **17**:e0265823. doi:[10.1371/journal.pone.0265823](https://doi.org/10.1371/journal.pone.0265823)
- Du Bois WEB. 1900. The American Negro at Paris **XXII**:575–577.
- Ed Hawkins. 2018. [2018 Visualisation Update](#). *Climate Lab Book*.
- Elias M, Aufaure M-A, Bezerianos A. 2013. Storytelling in Visual Analytics Tools for Business Intelligence In: Hutchison D, Kanade T, Kittler J, Kleinberg JM, Mattern F, Mitchell JC, Naor M, Nierstrasz O, Pandu Rangan C, Steffen B, Sudan M, Terzopoulos D, Tygar D, Vardi MY, Weikum G, Kotzé P, Marsden G, Lindgaard G, Wesson J, Winckler M, editors. Human-Computer Interaction – INTERACT 2013. Berlin, Heidelberg: Springer Berlin Heidelberg. pp. 280–297. doi:[10.1007/978-3-642-40477-1\\_18](https://doi.org/10.1007/978-3-642-40477-1_18)
- Elliott M, Nothelfer C, Xiong C, Szafir D. 2020. [A Design Space of Vision Science Methods for Visualization Research](#). *arXiv:200906855 [cs]*.
- Elting LS, Martin CG, Cantor SB, Rubenstein EB. 1999. Influence of data display formats on physician investigators' decisions to stop clinical trials: Prospective trial with repeated measures. *BMJ* **318**:1527–1531. doi:[10.1136/bmj.318.7197.1527](https://doi.org/10.1136/bmj.318.7197.1527)
- Fan A, Ma Y, Mancenido M, Maciejewski R. 2022. Annotating Line Charts for Addressing Deception CHI Conference on Human Factors in Computing Systems. New Orleans LA USA: ACM. pp. 1–12. doi:[10.1145/3491102.3502138](https://doi.org/10.1145/3491102.3502138)
- Feldman-Stewart D, Kocovski N, McConnell BA, Brundage MD, Mackillop WJ. 2000. Perception of Quantitative Information for Treatment Decisions. *Medical Decision Making* **20**:228–238. doi:[10.1177/0272989X0002000208](https://doi.org/10.1177/0272989X0002000208)
- Forrest J. 2018. [W. E. B. Du Bois' staggering Data Visualizations are as powerful today as they were in 1900 \(Part 1\)](#). *Nightingale*.

- Franconeri SL, Padilla LM, Shah P, Zacks JM, Hullman J. 2021. The Science of Visual Data Communication: What Works. *Psychological Science in the Public Interest* **22**:110–161. doi:[10.1177/15291006211051956](https://doi.org/10.1177/15291006211051956)
- Freeman ALJ, Kerr J, Recchia G, Schneider CR, Lawrence ACE, Finikarides L, Luoni G, Dryhurst S, Spiegelhalter D. 2021. Communicating personalized risks from COVID-19: Guidelines from an empirical study. *Royal Society Open Science* **8**:rsos.201721, 201721. doi:[10.1098/rsos.201721](https://doi.org/10.1098/rsos.201721)
- Friel SN, Curcio FR, Bright GW. 2001. Making Sense of Graphs: Critical Factors Influencing Comprehension and Instructional Implications. *Journal for Research in Mathematics Education* **32**:124. doi:[10.2307/749671](https://doi.org/10.2307/749671)
- Friendly M. 2008. [The Golden Age of Statistical Graphics](#). *Statistical Science* **23**:502–535.
- Friendly M. 2006. A Brief History of Data Visualization,.
- Friendly M, Andrews R. 2021. The radiant diagrams of Florence Nightingale. *SORT Statistics and Operations Research Transactions* 3–18. doi:[10.2436/20.8080.02.106](https://doi.org/10.2436/20.8080.02.106)
- Galesic M, Garcia-Retamero R. 2011. Graph Literacy: A Cross-Cultural Comparison. *Medical Decision Making* **31**:444–457. doi:[10.1177/0272989X10373805](https://doi.org/10.1177/0272989X10373805)
- Garcia-Retamero R, Cokely ET, Ghazal S, Joeris A. 2016. Measuring Graph Literacy without a Test: A Brief Subjective Assessment. *Medical Decision Making* **36**:854–867. doi:[10.1177/0272989X16655334](https://doi.org/10.1177/0272989X16655334)
- Ge LW, Cui Y, Kay M. 2023. CALVI: Critical Thinking Assessment for Literacy in VisualizationsProceedings of the 2023 CHI Conference on Human Factors in Computing Systems. Hamburg Germany: ACM. pp. 1–18. doi:[10.1145/3544548.3581406](https://doi.org/10.1145/3544548.3581406)
- Gillan DJ, Richman EH. 1994. Minimalism and the Syntax of Graphs. *Human Factors: The Journal of the Human Factors and Ergonomics Society* **36**:619–644. doi:[10.1177/001872089403600405](https://doi.org/10.1177/001872089403600405)
- Gleicher M, Correll M, Nothelfer C, Franconeri S. 2013. Perception of Average Value in Multiclass Scatterplots. *IEEE Transactions on Visualization and Computer Graphics* **19**:2316–2325. doi:[10.1109/TVCG.2013.183](https://doi.org/10.1109/TVCG.2013.183)
- Grice P. 1975. Logic and Conversation In: Cole P, Morgan JL, editors. Syntax and Semantics Vol.3: Speech Acts. New York: Academic Press. pp. 41–58.
- Hans Rosling. 2006. [The Best Stats You've Ever Seen](#).
- Harrison L, Yang F, Franconeri S, Chang R. 2014. Ranking Visualizations of Correlation

- Using Weber's Law. *IEEE Transactions on Visualization and Computer Graphics* **20**:1943–1952. doi:[10.1109/TVCG.2014.2346979](https://doi.org/10.1109/TVCG.2014.2346979)
- Heer J, Bostock M. 2010. Crowdsourcing graphical perception: Using mechanical turk to assess visualization design 10.
- Hollands JG, Spence I. 1998. Judging proportion with graphs: The summation model. *Applied Cognitive Psychology* **12**:173–190. doi:[10.1002/\(SICI\)1099-0720\(199804\)12:2<173::AID-ACP499>3.0.CO;2-K](https://doi.org/10.1002/(SICI)1099-0720(199804)12:2<173::AID-ACP499>3.0.CO;2-K)
- Hong M-H, Witt JK, Szafir DA. 2021. The Weighted Average Illusion: Biases in Perceived Mean Position in Scatterplots. *IEEE Transactions on Visualization and Computer Graphics* 1–1. doi:[10.1109/TVCG.2021.3114783](https://doi.org/10.1109/TVCG.2021.3114783)
- Huff D. 1993. How to lie with statistics. New York: Norton.
- Kale A, Kay M, Hullman J. 2020. [Visual Reasoning Strategies and Satisficing: How Uncertainty Visualization Design Impacts Effect Size Judgments and Decisions.](#) *arXiv:200714516 [cs]*.
- Kennedy H, Hill RL, Aiello G, Allen W. 2016. The work that visualisation conventions do. *Information, Communication & Society* **19**:715–735. doi:[10.1080/1369118X.2016.1153126](https://doi.org/10.1080/1369118X.2016.1153126)
- Kerns SH, Wilmer JB. 2021. Two Graphs Walk into a Bar: Readout-based Measurement Reveals the Bar-tip Limit Error, a Common, Categorical Misinterpretation of Bar Graphs. *Journal of Vision* 35.
- Kindlmann G, Scheidegger C. 2014. An Algebraic Process for Visualization Design. *IEEE Transactions on Visualization and Computer Graphics* **20**:2181–2190. doi:[10.1109/TVCG.2014.2346325](https://doi.org/10.1109/TVCG.2014.2346325)
- Kintisch E. 2019. New climate “stripes” reveal how much hotter your hometown has gotten in the past century. *Science*. doi:[10.1126/science.aay5354](https://doi.org/10.1126/science.aay5354)
- Kosara R. 2016. An Empire Built On Sand: Reexamining What We Think We Know About Visualization Proceedings of the Beyond Time and Errors on Novel Evaluation Methods for Visualization - BELIV '16. Baltimore, MD, USA: ACM Press. pp. 162–168. doi:[10.1145/2993901.2993909](https://doi.org/10.1145/2993901.2993909)
- Kosara R, Mackinlay J. 2013. Storytelling: The Next Step for Visualization. *Computer* **46**:44–50. doi:[10.1109/MC.2013.36](https://doi.org/10.1109/MC.2013.36)
- Kosslyn SM. 1985. Graphics and Human Information Processing: A Review of Five Books. *Current Biology* **25**:R134–R136. doi:[10.1016/j.cub.2014.12.018](https://doi.org/10.1016/j.cub.2014.12.018)
- Lam H, Bertini E, Isenberg P, Plaisant C, Carpendale S. 2012. Empirical Studies in



- Information Visualization: Seven Scenarios. *IEEE Transactions on Visualization and Computer Graphics* **18**:1520–1536. doi:[10.1109/TVCG.2011.279](https://doi.org/10.1109/TVCG.2011.279)
- Lauer C, O'Brien S. 2020. The Deceptive Potential of Common Design Tactics Used in Data Visualizations. Proceedings of the 38th ACM International Conference on Design of Communication. Denton TX USA: ACM. pp. 1–9. doi:[10.1145/3380851.3416762](https://doi.org/10.1145/3380851.3416762)
- Lee S, Kwon BC, Yang J, Lee BC, Kim S-H. 2019. The Correlation between Users' Cognitive Characteristics and Visualization Literacy. *Applied Sciences* **9**:488. doi:[10.3390/app9030488](https://doi.org/10.3390/app9030488)
- Lem S, Onghena P, Verschaffel L, Van Dooren W. 2014. [Interpreting histograms. As easy as it seems?](#) *European Journal of Psychology of Education* **29**:557–575.
- Lin S, Fortuna J, Kulkarni C, Stone M, Heer J. 2013. Selecting Semantically-Resonant Colors for Data Visualization. *Computer Graphics Forum* **32**:401–410. doi:[10.1111/cgf.12127](https://doi.org/10.1111/cgf.12127)
- Lohse GL. 1993. A Cognitive Model for Understanding Graphical Perception 37.
- MacGregor N. 2010. A history of the world in 100 objects. London: Allen Lane.
- Mackinlay J. 1986. Automating the design of graphical presentations of relational information. *ACM Transactions on Graphics* **5**:110–141. doi:[10.1145/22949.22950](https://doi.org/10.1145/22949.22950)
- Magnello ME. 2012. Victorian statistical graphics and the iconography of Florence Nightingale's polar area graph. *BSHM Bulletin: Journal of the British Society for the History of Mathematics* **27**:13–37. doi:[10.1080/17498430.2012.618102](https://doi.org/10.1080/17498430.2012.618102)
- Mantri P, Subramonyam H, Michal AL, Xiong C. 2022. How Do Viewers Synthesize Conflicting Information from Data Visualizations? *IEEE Transactions on Visualization and Computer Graphics* 1–11. doi:[10.1109/TVCG.2022.3209467](https://doi.org/10.1109/TVCG.2022.3209467)
- Mendonça A, Delazari L. 2014. Testing Subjective Preference and Map Use Performance: Use of Web Maps for Decision Making in the Public Health Sector. *Cartographica: The International Journal for Geographic Information and Geovisualization* **49**:114–126. doi:[10.3138/carto.49.2.1455](https://doi.org/10.3138/carto.49.2.1455)
- Newman GE, Scholl BJ. 2012. Bar graphs depicting averages are perceptually misinterpreted: The within-the-bar bias. *Psychonomic Bulletin & Review* **19**:601–607. doi:[10.3758/s13423-012-0247-5](https://doi.org/10.3758/s13423-012-0247-5)
- Nothelfer C, Franconeri S. 2020. Measures of the Benefit of Direct Encoding of Data Deltas for Data Pair Relation Perception. *IEEE Transactions on Visualization and Computer Graphics* **26**:311–320. doi:[10.1109/TVCG.2019.2934801](https://doi.org/10.1109/TVCG.2019.2934801)

- Nothelfer C, Gleicher M, Franconeri S. 2017. Redundant encoding strengthens segmentation and grouping in visual displays of data. *Journal of Experimental Psychology: Human Perception and Performance* **43**:1667–1676. doi:[10.1037/xhp0000314](https://doi.org/10.1037/xhp0000314)
- Obie HO, Chua C, Avazpour I, Abdelrazek M, Grundy J, Bednarz T. 2019. A study of the effects of narration on comprehension and memorability of visualisations. *Journal of Computer Languages* **52**:113–124. doi:[10.1016/j.col.2019.04.006](https://doi.org/10.1016/j.col.2019.04.006)
- Okan Y, Galesic M, Garcia-Retamero R. 2016. How People with Low and High Graph Literacy Process Health Graphs: Evidence from Eye-tracking: Graph Literacy and Health Graph Processing. *Journal of Behavioral Decision Making* **29**:271–294. doi:[10.1002/bdm.1891](https://doi.org/10.1002/bdm.1891)
- Okan Y, Garcia-Retamero R, Cokely ET, Maldonado A. 2012a. Individual Differences in Graph Literacy: Overcoming Denominator Neglect in Risk Comprehension. *Journal of Behavioral Decision Making* **25**:390–401. doi:[10.1002/bdm.751](https://doi.org/10.1002/bdm.751)
- Okan Y, Garcia-Retamero R, Galesic M, Cokely ET. 2012b. When Higher Bars Are Not Larger Quantities: On Individual Differences in the Use of Spatial Information in Graph Comprehension. *Spatial Cognition and Computation* **12**:195–218. doi:[10.1080/13875868.2012.659302](https://doi.org/10.1080/13875868.2012.659302)
- Okan Y, Janssen E, Galesic M, Waters EA. 2019. Using the Short Graph Literacy Scale to Predict Precursors of Health Behavior Change. *Medical Decision Making* **39**:183–195. doi:[10.1177/0272989X19829728](https://doi.org/10.1177/0272989X19829728)
- Okan Y, Stone ER, Bruine De Bruin W. 2018. Designing Graphs that Promote Both Risk Understanding and Behavior Change: Graphs Promoting Risk Understanding and Behavior Change. *Risk Analysis* **38**:929–946. doi:[10.1111/risa.12895](https://doi.org/10.1111/risa.12895)
- Padilla LM, Creem-Regehr SH, Hegarty M, Stefanucci JK. 2018. Decision making with visualizations: A cognitive framework across disciplines. *Cognitive Research: Principles and Implications* **3**:29. doi:[10.1186/s41235-018-0120-9](https://doi.org/10.1186/s41235-018-0120-9)
- Pandey AV, Rall K, Satterthwaite ML, Nov O, Bertini E. 2015. How Deceptive are Deceptive Visualizations?: An Empirical Analysis of Common Distortion Techniques. Proceedings of the 33rd Annual ACM Conference on Human Factors in Computing Systems - CHI '15. Seoul, Republic of Korea: ACM Press. pp. 1469–1478. doi:[10.1145/2702123.2702608](https://doi.org/10.1145/2702123.2702608)
- Pentoney CS, Berger DE. 2016. Confidence Intervals and the Within-the-Bar Bias. *The American Statistician* **70**:215–220. doi:[10.1080/00031305.2016.1141706](https://doi.org/10.1080/00031305.2016.1141706)

- Pinker S. 1990. A theory of graph comprehension. *Artificial Intelligence and the Future of Testing*. pp. 73–126.
- Reda K, Szafir DA. 2021. Rainbows Revisited: Modeling Effective Colormap Design for Graphical Inference. *IEEE Transactions on Visualization and Computer Graphics* **27**:1032–1042. doi:[10.1109/TVCG.2020.3030439](https://doi.org/10.1109/TVCG.2020.3030439)
- Rensink R. 2021. How Not to Study a Visualization 4.
- Reyna VF. 2008. A Theory of Medical Decision Making and Health: Fuzzy Trace Theory. *Medical Decision Making* **28**:850–865. doi:[10.1177/0272989X08327066](https://doi.org/10.1177/0272989X08327066)
- Reyna VF, Brainerd CJ. 1991. Fuzzy-trace theory and framing effects in choice: Gist extraction, truncation, and conversion. *Journal of Behavioral Decision Making* **4**:249–262. doi:[10.1002/bdm.3960040403](https://doi.org/10.1002/bdm.3960040403)
- Robertson G, Fernandez R, Fisher D, Lee B, Stasko J. 2008. Effectiveness of Animation in Trend Visualization. *IEEE Transactions on Visualization and Computer Graphics* **14**:1325–1332. doi:[10.1109/TVCG.2008.125](https://doi.org/10.1109/TVCG.2008.125)
- Romano A, Sotis C, Dominioni G, Guidi S. 2020. The scale of COVID-19 graphs affects understanding, attitudes, and policy preferences. *Health Economics* **29**:1482–1494. doi:[10.1002/hec.4143](https://doi.org/10.1002/hec.4143)
- Saket B, Endert A, Demiralp Ç. 2019. Task-Based Effectiveness of Basic Visualizations. *IEEE Transactions on Visualization and Computer Graphics* **25**:2505–2512. doi:[10.1109/TVCG.2018.2829750](https://doi.org/10.1109/TVCG.2018.2829750)
- Sandman PM, Weinstein ND, Miller P. 1994. High Risk or Low: How Location on a "Risk Ladder" Affects Perceived Risk. *Risk Analysis* **14**:35–45. doi:[10.1111/j.1539-6924.1994.tb00026.x](https://doi.org/10.1111/j.1539-6924.1994.tb00026.x)
- Satyanarayan A, Moritz D, Wongsuphasawat K, Heer J. 2017. Vega-Lite: A Grammar of Interactive Graphics. *IEEE Transactions on Visualization and Computer Graphics* **23**:341–350. doi:[10.1109/TVCG.2016.2599030](https://doi.org/10.1109/TVCG.2016.2599030)
- Shah RC, Kesan JP. 2006. Policy through software defaults. *Proceedings of the 2006 National Conference on Digital Government Research - Dg.o '06*. San Diego, California: ACM Press. p. 265. doi:[10.1145/1146598.1146670](https://doi.org/10.1145/1146598.1146670)
- Simkin D, Hastie R. 1987. An Information-Processing Analysis of Graph Perception. *Journal of the American Statistical Association* **82**:454–465. doi:[10.2307/2289447](https://doi.org/10.2307/2289447)
- Spence I, Lewandowsky S. 1991. Displaying proportions and percentages. *Applied Cognitive Psychology* **5**:61–77. doi:[10.1002/acp.2350050106](https://doi.org/10.1002/acp.2350050106)

- Stock WA, Behrens JT. 1991. Box, Line, and Midgap Plots: Effects of Display Characteristics on the Accuracy and Bias of Estimates of Whisker Length.
- Stolte C, Tang D, Hanrahan P. 2002. Polaris: A system for query, analysis, and visualization of multidimensional relational databases. *IEEE Transactions on Visualization and Computer Graphics* **8**:52–65. doi:[10.1109/2945.981851](https://doi.org/10.1109/2945.981851)
- Stone ER, Bruin W, Wilkins AM, Boker EM, MacDonald Gibson J. 2017. Designing Graphs to Communicate Risks: Understanding How the Choice of Graphical Format Influences Decision Making. *Risk Analysis* **37**:612–628. doi:[10.1111/risa.12660](https://doi.org/10.1111/risa.12660)
- Stone ER, Gabard AR, Groves AE, Lipkus IM. 2015. Effects of Numerical Versus Foreground-Only Icon Displays on Understanding of Risk Magnitudes. *Journal of Health Communication* **20**:1230–1241. doi:[10.1080/10810730.2015.1018594](https://doi.org/10.1080/10810730.2015.1018594)
- Stone ER, Sieck WR, Bull BE, Frank Yates J, Parks SC, Rush CJ. 2003. Foreground:background salience: Explaining the effects of graphical displays on risk avoidance. *Organizational Behavior and Human Decision Processes* **90**:19–36. doi:[10.1016/S0749-5978\(03\)00003-7](https://doi.org/10.1016/S0749-5978(03)00003-7)
- Strain G, Stewart AJ, Warren P, Jay C. 2023. The Effects of Contrast on Correlation Perception in Scatterplots. *International Journal of Human-Computer Studies* **176**:103040. doi:[10.1016/j.ijhcs.2023.103040](https://doi.org/10.1016/j.ijhcs.2023.103040)
- Suh A, Mosca A, Robinson S, Pham Q, Cashman D, Ottley A, Chang R. 2022. Inferential Tasks as an Evaluation Technique for Visualization 5 pages. doi:[10.2312/EVS.20221086](https://doi.org/10.2312/EVS.20221086)
- Szafir DA. 2018. Modeling Color Difference for Visualization Design. *IEEE Transactions on Visualization and Computer Graphics* **24**:392–401. doi:[10.1109/TVCG.2017.2744359](https://doi.org/10.1109/TVCG.2017.2744359)
- Taylor BG, Anderson LK. 1986. Misleading Graphs: Guidelines for the Accountant. *Journal of Accountancy* **162**:126–135.
- Tufte ER. 1986. The visual display of quantitative information. *Information Design Journal* **4**:235–236. doi:[10.1075/idj.4.3.12cos](https://doi.org/10.1075/idj.4.3.12cos)
- Tversky A, Kahneman D. 1981. The Framing of Decisions and the Psychology of Choice. *Science* **211**:453–458. doi:[10.1126/science.7455683](https://doi.org/10.1126/science.7455683)
- Tversky A, Kahneman D. 1974. Judgment under Uncertainty: Heuristics and Biases. *Science* **185**:1124–1131. doi:[10.1126/science.185.4157.1124](https://doi.org/10.1126/science.185.4157.1124)
- Tversky B. 1997. Cognitive Principles of Graphic Displays Proceedings of the AAAI 1997 Fall Symposium on Reasoning with Diagrammatic Representations. Menlo Park, CA: AAAI Press,. pp. 116–124.

- Wainer H. 1984. How to Display Data Badly. *The American Statistician* **38**:137. doi:[10.2307/2683253](https://doi.org/10.2307/2683253)
- Ware C, Stone M, Szafir DA. 2023. Rainbow Colormaps Are Not All Bad. *IEEE Computer Graphics and Applications* **43**:88–93. doi:[10.1109/MCG.2023.3246111](https://doi.org/10.1109/MCG.2023.3246111)
- Wickham H. 2011. ggplot2: ggplot2. *Wiley Interdisciplinary Reviews: Computational Statistics* **3**:180–185. doi:[10.1002/wics.147](https://doi.org/10.1002/wics.147)
- Wickham H. 2010. A Layered Grammar of Graphics. *Journal of Computational and Graphical Statistics* **19**:3–28. doi:[10.1198/jcgs.2009.07098](https://doi.org/10.1198/jcgs.2009.07098)
- Wilkinson L. 2005. The grammar of graphics, 2nd ed. ed, Statistics and computing. New York: Springer.
- Witt JK. 2019. Graph Construction: An Empirical Investigation on Setting the Range of the Y-Axis. *Meta-Psychology* **2**:1–20. doi:[10.15626/MP.2018.895](https://doi.org/10.15626/MP.2018.895)
- Woodin G, Winter B, Padilla L. 2022. Conceptual Metaphor and Graphical Convention Influence the Interpretation of Line Graphs. *IEEE Transactions on Visualization and Computer Graphics* **28**:1209–1221. doi:[10.1109/TVCG.2021.3088343](https://doi.org/10.1109/TVCG.2021.3088343)
- Xiong C, Lee-Robbins E, Zhang I, Gaba A, Franconeri S. 2022. Reasoning Affordances with Tables and Bar Charts. *IEEE Transactions on Visualization and Computer Graphics* 1–13. doi:[10.1109/TVCG.2022.3232959](https://doi.org/10.1109/TVCG.2022.3232959)
- Xiong C, Setlur V, Bach B, Lin K, Koh E, Franconeri S. 2021. [Visual Arrangements of Bar Charts Influence Comparisons in Viewer Takeaways](https://arxiv.org/abs/2108.06370). *arXiv:210806370 [cs]*.
- Xiong C, Weelden L van, Franconeri S. 2019. The Curse of Knowledge in Visual Data Communication. *IEEE Transactions on Visualization and Computer Graphics* 1–1. doi:[10.1109/TVCG.2019.2917689](https://doi.org/10.1109/TVCG.2019.2917689)
- Yang BW, Vargas Restrepo C, Stanley ML, Marsh EJ. 2021. Truncating Bar Graphs Persistently Misleads Viewers. *Journal of Applied Research in Memory and Cognition* S2211368120300978. doi:[10.1016/j.jarmac.2020.10.002](https://doi.org/10.1016/j.jarmac.2020.10.002)
- Yang L, Xiong C, Wong JK, Wu A, Qu H. 2023. Explaining With Examples: Lessons Learned From Crowdsourced Introductory Description of Information Visualizations. *IEEE Transactions on Visualization and Computer Graphics* **29**:1638–1650. doi:[10.1109/TVCG.2021.3128157](https://doi.org/10.1109/TVCG.2021.3128157)
- Yuan L, Haroz S, Franconeri S. 2019. Perceptual proxies for extracting averages in data visualizations. *Psychonomic Bulletin & Review* **26**:669–676. doi:[10.3758/s13423-018-1525-7](https://doi.org/10.3758/s13423-018-1525-7)

- Zacks J, Levy E, Tversky B, Schiano DJ. 1998. Reading bar graphs: Effects of extraneous depth cues and graphical context. *Journal of Experimental Psychology: Applied* **4**:119–138. doi:[10.1037/1076-898X.4.2.119](https://doi.org/10.1037/1076-898X.4.2.119)
- Zacks J, Tversky B. 1999. Bars and lines: A study of graphic communication. *Memory & Cognition* **27**:1073–1079. doi:[10.3758/BF03201236](https://doi.org/10.3758/BF03201236)
- Ziemkiewicz C, Kosara R. 2008. The Shaping of Information by Visual Metaphors. *IEEE Transactions on Visualization and Computer Graphics* **14**:1269–1276. doi:[10.1109/TVCG.2008.171](https://doi.org/10.1109/TVCG.2008.171)

## Chapter 3

# Experimental, Analytical, and Computational Methodology

### 3.1 Introduction

The knowledge generated in a research project is necessarily shaped by the methods of inquiry. A recent survey of visualisation researchers revealed variation in conceptions of how progress is made in the field, with multiple approaches for generating knowledge (Correll, 2022). In this chapter, I discuss the epistemological approach which underlies this thesis. This provides a backdrop to the subsequent empirical work and a justification for my choices. It is also necessary to acknowledge that this methodology is one of many possible methodologies, and each carries its own implications. By explicitly discussing my decisions, I recognise that they inevitably influence my findings. This reflects the fact that an epistemological approach imposes a particular perspective, unavoidably generating a somewhat narrow view on the topic of interest.

### 3.2 Experimental Methodology

Selecting a research method involves considering the most suitable type of data for addressing the research question. To understand how data visualisation design choices affect interpretations of absolute magnitude, testing hypotheses using controlled experiments is highly appropriate. This allows for systematic measurement of viewers' judgements and isolates the graphical features of interest from extraneous features. Controlling for the influence of other variables helps establish a causal link between a manipulation and

cognition (Barbosa et al., 2021). Experimental methods are well-established in visualisation research for generating robust empirical evidence on the effects of design choices (Abdul-Rahman et al., 2019).

### **3.2.1 Experimental Design and Ecological Validity**

The purpose of psychological studies on data visualisation is to develop an understanding of viewers' interpretations. However, latent variables cannot be interrogated directly, and must be 'operationalised' to enable analysis. For example, interpretations of the magnitude of numerical values must be captured through measurable responses which correspond to underlying mental representations. Thus, experimental methods rely on the use of dependent variables which faithfully reflect actual cognition. Designing experiments also requires compromise between ecological validity (the degree to which the experiment reflects a realistic scenario) and experimental control (the degree to which the researcher dictates aspects of the experiment). In this thesis, I have strived for realism where possible, but have prioritised experimental control in order to ensure the robustness of findings. In visualisation research in particular, it is often necessary to control for differences in participants' knowledge by presenting artificial or abstract data (Lam et al., 2012). Whilst qualitative studies (e.g., in-depth interviews), may produce richer data than experiments, they do not provide the precision required to systematically evaluate biases in interpretation.

This is a largely positivist approach, concerned with verifiable results which can be generalised beyond the experiment to describe a cognitive mechanism. However, there is also arguably a *postmodern* quality to highly controlled experiments (Mayrhofer et al., 2021). That is, a controlled experiment can be considered a constructed, stimulated setting, with contrived tasks and stimuli that do not precisely reflect the 'reality' under investigation (i.e., spontaneous judgements of authentic data visualisations). Recognising this does not invalidate conclusions from experimental studies, but requires that generalisation of results is treated with caution.

### **3.2.2 Sample Size and Generalisability**

All experiments in this thesis were conducted online using Prolific.co, a website for recruiting research participants. This provided access to a diverse group of participants, which contrasts with the relative homogeneity of a student population, typically employed in



experimental psychology research. Furthermore, online experiments provide the ability to easily collect data from a large number of participants, which reduces the chance of generating false positives during analysis. In addition to using large participant samples, employing multiple trials per condition helps establish robust effects which are not vulnerable to the particular characteristics a single trial. Similarly, generating generalisable knowledge about mental processing of visualisations often requires multiple experiments. A single experiment is typically not sufficient for an understanding of cognitive factors in interpretation (Chen et al., 2020).

### 3.3 Analytical Methodology

#### 3.3.1 Linear Mixed-Effects Models

Large quantitative datasets from controlled experiments require appropriate statistical analysis. Determining whether an experimental manipulation has affected participants' interpretations involves examining variability between different experimental conditions against the background other variability in the dataset. Linear mixed-effects modelling offers a powerful and reliable approach to this task, and is used throughout this thesis. Much like my experimental methodology, this reflects a positivist approach, wherein conclusions are supported by mathematical verification.

#### 3.3.2 Fixed Effects and Random Effects

A central aspect of mixed-effects modelling is the distinction between fixed effects and random effects. Independent variables of interest are modelled as *fixed effects*; their influence on the dependent variable is the primary focus of an analysis. Other sources of variability are modelled as *random effects* in order to generate a more comprehensive model of a dataset. For a variable manipulated in an experiment, each relevant level is present in the dataset, so this variable should be modelled as a fixed effect. However, typically, only a *sample* of all possible participants or all possible stimuli are present in a dataset, so these variables should be modelled as random effects. At a minimum, specifying a random effect involves modelling *intercepts*: the average response at each level of the random effect (i.e., a separate baseline for each individual participant or experimental stimulus). Additionally, researchers may model *slopes*: the effect of the independent variable at each level of the random effect (i.e., the difference between conditions for each individual

participant or experimental stimulus). Thus, modelling random intercepts *and* random slopes attempts to capture more variability in a dataset than modelling random intercepts alone.

### 3.3.3 Benefits of Mixed-Effects Models

Modelling random effects is beneficial as a means of testing the *generalisability* of an fixed effect: whether it is robust when differences across participants and experimental stimuli are taken into account. For example, including random effects for experimental stimuli can improve prediction of whether the results will replicate when different stimuli are used (Judd et al., 2017). Furthermore, it is appropriate to recognise the dependencies between data points associated with the same participant, or the same experimental item, rather than incorrectly treating them as independent observations. Compared to simpler models, better parameter estimates in mixed-effects models decrease the likelihood of generating false positives (Singmann and Kellen, 2019).

### 3.3.4 Approaches to Model Construction

Barr et al. (2013) argue that researchers should construct ‘maximal’ models which reflect the full complexity of their experimental design. Therefore, in a fully-crossed design where there are observations at each level of the fixed effect for each participant and each experimental stimulus, researchers should employ random intercepts and slopes for participants and experimental stimuli. Barr et al. (2013) suggest that modelling all possible random effects increases statistical power without inflating Type 1 error. However, Bates et al. (2018) argue for a different approach, which acknowledges that building complex mixed-effects models is a complex process, and not all datasets are sufficiently rich to support such computations. Estimating a large number of parameters using a small number of observations can result in ‘overfitting’: the resulting estimates are not always reliable. Thus, maximal model structures can be over-ambitious. Including all possible random effects terms does not necessarily improve the modelling of fixed effects.

### 3.3.5 An Automated Approach

Whereas Barr et al.’s (2013) recommendations for specifying random effects structures are primarily informed by the experimental design, Bates et al.’s (2018) recommendations are primarily informed by the dataset. The approach to model construction used in this

thesis is influenced by both positions, and attempts to balance simplicity and explanatory power. An additional constraint is model convergence, which refers to the process of generating a solution when building a model. In this thesis, models are constructed using a two-stage process which is automated using the *buildmer* package in R (Voeten, 2022). First, this software attempts to build the maximal model and identify the most complex random effects structure which results in successful model convergence. Second, the software simplifies the model structure by removing random effects terms which do not contribute significantly to explaining variance in the dataset. This seeks to maximise the variability captured by the model whilst minimising the *redundant* random effects terms which may result in unreliable parameter estimates. Furthermore, using a computational process provides a consistent, rigorous, and transparent approach to model construction. A reproducible account of the steps preceding identification of each statistical model reduces the chance of human error and documents the process as well as its outcome (Rule et al., 2019).

### 3.3.6 Reporting Analyses

In addition to the construction of robust statistical models, the *reporting* of statistical analyses is another important consideration. Statistically significant results concerning differences between experimental conditions do not indicate how *substantial* differences are, so I report effect sizes in addition to p-values and test statistics (Wilkinson and Task Force on Statistical Inference, 1999). Following recommendations for mixed-effects modelling, I also report model structures alongside results, for transparency (Meteyard and Davies, 2020).

## 3.4 Reproducibility

In recent years, the typical model for conducting and publishing scientific research has been intensely scrutinised. This has prompted serious concern about the degree to which reported findings can be trusted. For example, Ioannidis (2005) estimated that published research may consist of more falsehoods than true assertions. Researchers also report that in the field of Psychology, many studies are not equipped to generate reliable results (Fraleigh and Vazire, 2014) and the literature is afflicted with a high rate of false-positive findings (Simmons et al., 2011). A large-scale project performing replications of psychology experiments revealed that the evidence for many established conclusions was not as

strong as initially reported (Open Science Collaboration, 2015). A survey of over 1500 researchers found widespread perception that science was facing a ‘crisis’ (Baker, 2016). However, this recognition also has provoked concerted efforts to address these problems in research, through the Open Science movement (Crüwell et al., 2019).

Recommendations for improving scientific research focus variously on different aspects of the research lifecycle. Improving how studies are conducted, reported, and evaluated requires targeted solutions. For example, rigorous methods and statistical analysis facilitate researchers in generating valid conclusions. Other practices, such as openly sharing data and code, increase transparency, providing crucial insight into how these conclusions were generated (Munafò et al., 2017). Peng (2011) suggests that the ultimate test of scientific claims is *replication*. This involves independently repeating an entire empirical investigation, thus generating new data to assess its consistency with an existing finding. However, this is resource-intensive. A different, albeit less rigorous, approach to evaluating scientific claims involves using a project’s original data and code to validate reported findings. If this is possible, the work is *reproducible*. By reusing existing resources, this approach is simpler than conducting a replication study, yet still facilitates assessment of whether reported results are reliable. To assist in the evaluation of research reproducibility, researchers can make relevant resources available.

The empirical work presented in this thesis has been conducted with a focus on ensuring reproducibility. The remaining contents of this chapter will review published work on best practices for sharing code, data, and computational environments, and outline the approach to reproducibility employed in this thesis.

### 3.5 Sharing Code and Data

There are many convincing arguments for openly sharing code and data. Scientific approaches benefit from the capacity to thoroughly assess the credibility of published work (Klein et al., 2018) and can independently authenticate other researchers’ conclusions (Blischak et al., 2019). Thus, supporting third parties in reproducing research can increase perceptions of its robustness and reliability (Sandve et al., 2013). This can also facilitate identification of errors in analysis (Klein et al., 2018). In addition to these motivating factors, authors may even appreciate the advantages of reproducible practices more than their audience (Piccolo and Frampton, 2016). For example, these practices can

save time and effort (Sandve et al., 2013), and permanently sharing resources provides insurance against the loss of those resources (Klein et al., 2018).

A textual description of analysis in a manuscript typically presents an somewhat incomplete and imprecise account of the analytical process (Piccolo and Frampton, 2016). Sharing code helps detail the precise journey from the original dataset to inferential statistics (Klein et al., 2018), providing a comprehensive account of this process. In the past, the possibility of issues or inconsistencies arising from computer code was overlooked (Plesser, 2018). However, it is now widely recognised that a computational analysis pipeline can present opportunities for error. Making code openly available permits *independent* reproduction of all computational processes (Stodden et al., 2016). This, in turn, can engender trust, promote collaboration, and facilitate new applications (Jiménez et al., 2017). Each stage of processing should be included (Sandve et al., 2013) and any files produced using the analytical pipeline should be expendable, since reproducing them using the code supplied should be trivial (Marwick et al., 2018). For full transparency, data should be supplied in a raw, unprocessed form (White et al., 2013). Keeping raw data separate from other files ensures that the original file is not altered and the stages of processing are clear (Marwick et al., 2018). Other resources, such as stimuli and experiment scripts should also be shared alongside data and code (Klein et al., 2018).

The FAIR principles (Wilkinson et al., 2016) propose that data (and metadata) should be Findable (easily discovered), Accessible (easily obtained), Interoperable (easily integrated with other tools), and Reusable (easily employed beyond their original use). The FAIR principles are also relevant for other computational tools, such as analysis scripts (Lamprecht et al., 2020), and share similarities with Open Source Software, which does not place limits on who may examine, adapt and extend the underlying code (Jiménez et al., 2017).

When sharing resources, a researcher's choices can assist or inhibit re-use (Chen et al., 2019)\*. For example, using non-proprietary file types ensures that third parties can readily access resources (White et al., 2013). Rather than personal or institutional websites, independent providers (e.g., Open Science Framework) are recommended for depositing these resources (Chen et al., 2019; Klein et al., 2018). Effective documentation is also valuable. A 'codebook' or 'data dictionary' can be used to explain the contents of a data file (Klein et al., 2018), inline comments can be used to explain code (Rule et al.,

2019), and a README file can be used to cover elementary information such as setup instructions (Lee, 2018). Documentation can also provide details on data collection and known issues in resources (White et al., 2013). Finally, licences contribute to a research project's longevity, and provide a clear statement for third parties, ensuring that their use of resources is appropriate (Jiménez et al., 2017). Where possible, lenient licences should be employed to avoid unnecessary restrictions (White et al., 2013).

### 3.5.1 The Importance of Public Sharing

It is fallacious to assert that if authors consistently shared data and code *on request*, freely available access would be unnecessary. To begin with, research papers outlive their authors, and requests obviously cannot be fulfilled by an author after they die (Klein et al., 2018). Empirical research further demonstrates why it is valuable to share resources *publicly*. In a study of 204 papers from a journal which *required* authors to provide data and code on request, only 44% delivered on this promise (Stodden et al., 2018). Where research code is not publicly available, various issues preclude procurement. These include local storage failures, restrictive institutional licences, concern about potential use, and concern about labour involved in providing support (Collberg and Proebsting, 2016). Provision of data and code on request simply cannot be guaranteed, motivating calls for public sharing. However, in the field of data visualisation research, public sharing has historically been uncommon. Of papers submitted to the VIS 2017 conference, 15% shared materials openly and 6% shared data openly (Haroz, 2018). Greater transparency could help to increase the credibility of data visualisation research and also potentially facilitate identification and rectification of issues in published work (Kosara and Haroz, 2018).

Researchers' working practices and technological solutions both contribute to reproducibility. Whilst it has been suggested that behaviour and technology play *equal* roles (Sandve et al., 2013), others argue that innovations have been so effective that researchers' engagement with these tools is now the primary driver of reproducible practices (Grüning et al., 2018). Researchers report that several factors impede or deter their sharing of research data, including lack of expertise, lack of precedent, and lack of time (Houtkoop et al., 2018).

## 3.6 Effective Programming Practices

Conducting analysis using a programmatic approach has three main benefits over manual processing: increased reproducibility, increased efficiency, and reduced error (Sandve et al., 2013). Writing functions in a modular style can avoid redundant repetition, promote comprehension, and support reuse of code (Wilson et al., 2017). Similar recommendations include splitting code into appropriate chunks which each achieve a clearly-defined goal (Rule et al., 2019). These techniques share many similarities with the Unix philosophy, an approach to computer programming which emphasises simplicity, modularity, and reusability (Gancarz, 2003).

The task of preparing data prior to analysis is an important aspect of working with data. Wickham (2014) presents a set of tools and underlying theory for this task, arguing that analysis can be facilitated by ensuring that data is in an appropriate structure. The recommended structure, known as ‘tidy’ data, consists of a column for each variable (each type of measurement) and a row for each observation (each unit measured). A principled approach simplifies the process of creating a tidy dataset using Wickham’s functions. Because each function treats data in a standardised manner, various functions can be employed in concert. The collection of R packages containing these functions (‘the Tidyverse’) was designed with a concern for *humans*, not just computational performance (Wickham et al., 2019), so Tidyverse-style code is likely to promote comprehension (Bertin and Baumer, 2021).

Several other coding behaviours can facilitate reproducibility. For example, *absolute* file paths refer to a specific directory on a user’s machine, which will not be replicated on other users’ machines. Using *relative* file paths, which locate files in relation to the project directory, ensure code is *portable* and can be used on any machine (Bertin and Baumer, 2021). Additionally, third parties cannot independently verify findings if only an approximate resemblance is achieved. Therefore, for any process involving random number generation, a random seed should be specified within the script, to ensure exact reproduction of results (Sandve et al., 2013).

### 3.6.1 Literate Programming and Dynamic Documents

Knuth’s (1984) novel perspective on comprehensibility in computer programming has been influential in the literature on computational reproducibility. Knuth’s premise is

that a programming script should not be regarded primarily as a set of instructions for a computer to follow, but as a tool to assist humans in understanding those instructions. This approach, known as ‘literate programming’, involves pairing code with corresponding text, such that reporting and documentation are closely linked to underlying code (Piccolo and Frampton, 2016; Sandve et al., 2013). Dynamic documents allow authors to mix code and narrative within a single file, with the results updated whenever the document is rendered. Producing (and re-producing) an entire manuscript using a dynamic document offers opportunities to easily observe the implementation of code used for each aspect of analysis (Peikert and Brandmaier, 2021). In addition to descriptive and inferential statistics, data visualisations may also be rendered dynamically (FitzJohn et al., 2014). This efficient format enhances transparency (Holmes et al., 2021), supports interactivity (Rule et al., 2019) and avoids errors which can occur when manually collating results (Peikert and Brandmaier, 2021). Including computationally-expensive code (e.g., complex statistical models) within a dynamic document can be problematic since this code is executed every time the document is rendered (FitzJohn et al., 2014). However, capacity for model caching provides a convenient antidote. This facilitates access to results by storing the output from models, which is then only updated when relevant data and code are updated.

### 3.7 Computational Environments

Providing data and code is necessary, but not sufficient, for enabling reproducibility. For example, research has found that even when the nominally required resources are available, it is not always possible to reproduce results exactly (Stodden et al., 2018), or even to execute the code (Collberg and Proebsting, 2016). In a high-profile case, a publicly-accessible Python script for processing organic chemistry data relied on the ordering of files by the Windows operating system, producing erroneous results for Linux users (Bhandari Neupane et al., 2019). A study using an automated approach to test the execution of 379 Python scripts from academic research found that success depended in part on the Python version used and the presence of files capturing dependencies (Trisovic et al., 2021). Another study using a similar approach to test over 9000 R scripts found that approximately three in four scripts produced errors when executed (Trisovic et al., 2022). Implementing a code-cleaning algorithm reduced this number, but the majority (56%) still failed to run successfully. This indicates that good programming practices can improve



code but cannot totally eliminate issues. A remaining source of error was incompatibility between R software versions and required packages. These studies illustrate that inability to recreate the *computational environment* used when originally running a script can prevent successful execution.

Peng (2011) argues that reproducibility can be characterised as a spectrum. Sharing code offers some benefits over a standalone publication, providing data increases reproducibility further, but ensuring that the code can be precisely executed is even better. Each researcher's unique preferences and proficiencies result in roughly the same number of computational environments as individual researchers, illustrating the utility of recording one's computational environment (Nüst et al., 2017). Additionally, software under continuous development, such as the Tidyverse collection of packages in R, is frequently updated, which means code can stop functioning unless specific versions are recorded (Holmes et al., 2021). Other software dependencies and parameter settings also complicate reproduction, requiring precision and comprehensiveness in documentation in order to achieve full *computational reproducibility* (Piccolo and Frampton, 2016).

### 3.7.1 Capturing Computational Environments Using Containers

Like many other aspects of reproducibility, innovations in software have made it possible for researchers to easily capture their computational environments. R package managers, such as *renv* (Ushey and Wickham, 2023) conveniently load specific package versions for individual projects. However, they do not guarantee computational reproducibility, because they do not preserve the version of R in the same way (Holmes et al., 2021) or support additional dependencies (Nüst et al., 2017; Peikert and Brandmaier, 2021). Containerisation technology offers an effective solution. A 'container' can capture a much greater extent of the computational environment than a package manager (Grüning et al., 2018). This technology also provides an efficient and principled approach for recreating the environment, compared to a list of instructions for manual execution (Marwick et al., 2018).

Docker (Merkel, 2014) is a popular tool for generating containers. This process begins with a Dockerfile: a text-based file which provides instructions for installing specific package versions and loading other dependencies and resources. The Dockerfile is used to build a Docker image, which captures the computational environment. When this image

is running, the environment is activated, and users may interact with this environment (Boettiger and Eddelbuettel, 2017; Nüst et al., 2020b).

Collating all dependency information in a single Dockerfile provides simplicity, and ensures that the original computational environment can be reproduced even after updating the software on a local machine. Since the primary objective is ensuring reproducibility, this approach prioritises openness and human readability over optimising performance [nust\_ten\_2020; Boettiger (2015)]. As well as simple implementations, complex arrangements can be accommodated, but present additional challenges. For example, dynamic document generation may also require specifying LaTeX dependencies (Boettiger, 2015).

### **3.7.2 Rocker for Capturing R Environments**

Researchers can save time and ensure consistency by using pre-existing Docker images (Nüst et al., 2020b). One particularly valuable example of this is Rocker (Boettiger and Eddelbuettel, 2014) which captures R environments for use in Docker. This tool provides portable R environments for use on any operating system, facilitating computational reproducibility. Consequently, any researcher can execute, edit, and extend R code in a replica of the environment originally used for its development (Boettiger and Eddelbuettel, 2014). Developing Rocker images involves a trade-off between generalisability and specificity. Images designed to be too widely applicable would be cumbersome, but images with overly-specific use cases would be hard to find (Boettiger and Eddelbuettel, 2017). The solution involves providing base images that are easily expanded for specific requirements, with various Rocker images ‘stacked’ together as required, avoiding unnecessary complexity (Nüst et al., 2020a).

### **3.7.3 Comparing Containers with Virtual Machines**

Virtual machines perform a similar function to containers. However a notable difference is that virtual machines are large, whilst containers are comparatively lightweight (Piccolo and Frampton, 2016). This difference is due to the fact that virtual machines use their own kernel, whereas containers use the operating system kernel provided by the local machine. This reduces the relative size of a container, and enhances its computational power (Cito et al., 2016). Thus, virtual machines may be considered more comprehensive than containers, offering a greater degree of separation from the characteristics of the host machine (Grüning et al., 2018; Piccolo and Frampton, 2016). However, containers are

typically compatible with version control systems (Piccolo and Frampton, 2016) and offer greater transparency (Nüst et al., 2020b). Furthermore, due to their modular features, making minor adaptations is trivial when using a container but comparatively prolonged when using a virtual machine.

### 3.8 Pragmatism Over Perfectionism

Despite the myriad recommendations for best practice, a principle often endorsed in the literature on reproducibility concerns the merits of small efforts. Taking some steps to increase reproducibility can still enhance a project's quality relative to overlooking this aspect altogether (Piccolo and Frampton, 2016). Withholding resources in pursuit of continuous refinement risks never sharing them at all. This fallacy is captured by the maxim 'the best is the enemy of the good'. Analysis code does not need to be perfect in order to be useful to others (Klein et al., 2018), and it is not possible to benefit from external inquiry if the code is not shared (Barnes, 2010). Barnes (2010) argues that perceived limitations simply reflect that the code works only for the specific scenario at hand; inessential improvements are by definition not required for basic functioning. Researchers are encouraged to accept these limitations and share their code anyway. In addition to code, this notion has also been applied to metadata (White et al., 2013) and containerisation (Nüst et al., 2020b).

### 3.9 The Approach to Reproducibility in This Thesis

The following describes the different aspects of reproducibility present in the subsequent empirical experiments in this thesis. Whilst this work does not follow a pre-defined workflow, the approach closely resembles workflows described in published work (e.g., Lissa et al., 2020; Peikert and Brandmaier, 2021).

#### 3.9.1 Data, Code, and Dynamic Documents

For each experiment in this thesis, raw data is provided. The only pre-processing of this data was the essential removal of sensitive information, which is transparently documented in corresponding scripts. All subsequent processing, from data cleaning to data wrangling, is included in a Quarto dynamic document (Allaire et al., 2022), which also includes all data analysis, visualisation, and accompanying text. Therefore, consistent with

the principles of literate programming, textual descriptions are presented in conjunction with corresponding code (Sandve et al., 2013). The Quarto document associated with each empirical chapter is openly available in its corresponding online repository, but each empirical chapter in this thesis consists of the *rendered* version of the document.

### 3.9.2 Docker Containers

Capturing software dependencies requires reproduction of the computational environment used (Boettiger, 2015). Each empirical chapter in this thesis is associated with a Dockerfile, which can be used to build a Docker container with the appropriate R version and package versions used during analysis. Employing Rocker images provides an Integrated Development Environment (RStudio), and speeds up construction of the Docker image. In each container, the entire chapter can be generated from scratch. The Dockerfiles also provide important project metadata in a human- and machine-readable format (Leipzig et al., 2021).

The following is the code from a Dockerfile used to reproduce the computational environment for the analysis conducted in Chapter 6:

First, I specify the Rocker image upon which the rest of the container will be built. This image includes R (version 4.2.1), plus the RStudio Integrated Development Environment, the Quarto publishing software, and the Tidyverse packages associated with this version of R.

```
FROM rocker/verse:4.2.1
```

Next, I add files to the image, including Quarto dynamic document itself, and its dependencies, including the bibliography associated with this chapter. These are mounted at the 'rstudio' directory.

```
ADD axis-extension.qmd /home/rstudio/
```

```
ADD _quarto.yml /home/rstudio/
```

```
ADD axis-extension.bib /home/rstudio/
```

Next, I add the contents of folders to the image. The first line adds the data files, the second line adds images displayed in the chapter, and the third line adds the cache containing the analysis models.

```
ADD data/ /home/rstudio/data/
```

```
ADD images/ /home/rstudio/images/
```

```
ADD axis-extension_cache/html/ /home/rstudio/axis-extension_cache/html
```

Next, I run two lines of R code to install and load the *renv* package. This package is a package manager which captures the versions of all other packages used in a project, and their dependencies, in a *renv.lock* file.

```
RUN R -e "install.packages('renv')"
```

```
RUN R -e "require(renv)"
```

Next, I make the *renv.lock* file available to Docker, so that it can access the specific package and dependency versions used for the original analysis.

```
COPY renv.lock renv.lock
```

Finally, I run a line of R code which installs the specific version of each package as specified in the *renv.lock* file.

```
RUN Rscript -e 'options(warn = 2); renv::restore(packages = c("ggribes",  
"buildmer", "broom.mixed", "lme4", "insight", "papaja", "magick",  
"patchwork", "ggpubr", "kableExtra", "emmeans", "knitr", "effectsize",  
"qwraps2", "report", "MuMIn", "shiny", "markdown"))'
```

### 3.9.3 Experiment Resources

In experimental psychology, sharing stimuli and experiment scripts is another important aspect of transparent research practice (Klein et al., 2018). All data visualisations shown to participants, along with all code used to generate those visualisations, has been made available. Experiments were programmed using PsychoPy, which developed as a tool for conducting open and reproducible research (Peirce et al., 2019). The underlying technology is open source, the experiment scripts use non-proprietary file formats, and the ability to specify particular software versions avoids new releases breaking older code. Its integration with GitLab version control software means that each experiment is packaged in a public online repository. An entire project's resources can be downloaded to a local machine, and an interactive version of the experiment can be run online.

### 3.10 Conclusion

In this chapter, I have outlined the experimental, analytical, and computational methodology used in this thesis. This involves conducting controlled experiments in order to systematically examine the cognitive mechanisms underlying the interpretation of data visualisations. Linear mixed-effects analysis of the resulting data provides a powerful inferential approach, with a model selection algorithm providing transparency, consistency, and statistical rigour. I have also discussed academic literature which explains how comprehensively sharing resources and embracing technological solutions can increase the credibility of published research. For each empirical chapter in this thesis, I share raw data alongside code packaged in a dynamic document, which illustrates exactly how the study's findings were generated. In addition, creating Docker containers for each study allows the analysis to be reproduced in its original computational environment. This comprehensive approach is uncommon in research on data visualisation, therefore this work may serve as an example of how research in this field can be made more transparent.

### References

- Abdul-Rahman A, Borgo R, Chen M, Edwards DJ, Fisher B. 2019. [Juxtaposing Controlled Empirical Studies in Visualization with Topic Developments in Psychology](#). *arXiv:190903786 [cs]*.
- Allaire JJ, Teague C, Xie Y, Dervieux C. 2022. Quarto. doi:[10.5281/ZENODO.5960048](#)
- Baker M. 2016. 1,500 scientists lift the lid on reproducibility. *Nature* **533**:452–454. doi:[10.1038/533452a](#)
- Barbosa J, Stein H, Summerfield C, Soto-Faraco S, Hyafil A. 2021. A practical guide for studying human behavior in the lab (preprint). PsyArXiv. doi:[10.31234/osf.io/tcmvp](#)
- Barnes N. 2010. Publish your computer code: It is good enough. *Nature* **467**:753–753. doi:[10.1038/467753a](#)
- Barr DJ, Levy R, Scheepers C, Tily HJ. 2013. Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language* **68**:255–278. doi:[10.1016/j.jml.2012.11.001](#)
- Bates D, Kliegl R, Vasishth S, Baayen H. 2018. [Parsimonious Mixed Models](#). *arXiv:150604967 [stat]*.
- Bertin AM, Baumer BS. 2021. Creating optimal conditions for reproducible data analysis in R with “fertile.” *Stat* **10**. doi:[10.1002/sta4.332](#)

- Bhandari Neupane J, Neupane RP, Luo Y, Yoshida WY, Sun R, Williams PG. 2019. Characterization of Leptazolines A–D, Polar Oxazolines from the Cyanobacterium *leptolyngbya* sp., Reveals a Glitch with the “Willoughby–Hoye” Scripts for Calculating NMR Chemical Shifts. *Organic Letters* **21**:8449–8453. doi:[10.1021/acs.orglett.9b03216](https://doi.org/10.1021/acs.orglett.9b03216)
- Blischak JD, Carbonetto P, Stephens M. 2019. Creating and sharing reproducible research code the workflowr way. *F1000Research* **8**:1749. doi:[10.12688/f1000research.20843.1](https://doi.org/10.12688/f1000research.20843.1)
- Boettiger C. 2015. An introduction to Docker for reproducible research, with examples from the R environment. *ACM SIGOPS Operating Systems Review* **49**:71–79. doi:[10.1145/2723872.2723882](https://doi.org/10.1145/2723872.2723882)
- Boettiger C, Eddelbuettel D. 2017. An Introduction to Rocker: Docker Containers for R. *The R Journal* **9**:527. doi:[10.32614/RJ-2017-065](https://doi.org/10.32614/RJ-2017-065)
- Boettiger C, Eddelbuettel D. 2014. Introducing Rocker: Docker for R (other). *Computer; information sciences*. doi:[10.59350/j1028-pq649](https://doi.org/10.59350/j1028-pq649)
- Chen M, Abdul-Rahman A, Laidlaw DH. 2020. [The Huge Variable Space in Empirical Studies for Visualization – A Challenge as well as an opportunity for Visualization Psychology](#).
- Chen X, Dallmeier-Tiessen S, Dasler R, Feger S, Fokianos P, Gonzalez JB, Hirvonsalo H, Kousidis D, Lavasa A, Mele S, Rodriguez DR, Šimko T, Smith T, Trisovic A, Trzcinska A, Tsanaktsidis I, Zimmermann M, Cranmer K, Heinrich L, Watts G, Hildreth M, Lloret Iglesias L, Lassila-Perini K, Neubert S. 2019. Open is not enough. *Nature Physics* **15**:113–119. doi:[10.1038/s41567-018-0342-2](https://doi.org/10.1038/s41567-018-0342-2)
- Cito J, Ferme V, Gall HC. 2016. Using Docker Containers to Improve Reproducibility in Software and Web Engineering Research In: Bozzon A, Cudre-Maroux P, Pautasso C, editors. *Web Engineering*. Cham: Springer International Publishing. pp. 609–612. doi:[10.1007/978-3-319-38791-8\\_58](https://doi.org/10.1007/978-3-319-38791-8_58)
- Collberg C, Proebsting TA. 2016. Repeatability in computer systems research. *Communications of the ACM* **59**:62–69. doi:[10.1145/2812803](https://doi.org/10.1145/2812803)
- Correll M. 2022. [Are We Making Progress In Visualization Research?](#)
- Crüwell S, Doorn J van, Etz A, Makel MC, Moshontz H, Niebaum JC, Orben A, Parsons S, Schulte-Mecklenbeck M. 2019. Seven Easy Steps to Open Science: An Annotated Reading List. *Zeitschrift für Psychologie* **227**:237–248. doi:[10.1027/2151-2604/a000387](https://doi.org/10.1027/2151-2604/a000387)

- FitzJohn R, Pennell M, Zanne A, Cornwell W. 2014. [Reproducible research is still a challenge](#).
- Fraley RC, Vazire S. 2014. The N-Pact Factor: Evaluating the Quality of Empirical Journals with Respect to Sample Size and Statistical Power. *PLoS ONE* **9**:e109019. doi:[10.1371/journal.pone.0109019](#)
- Gancarz M. 2003. Linux and the Unix philosophy. Amsterdam ; Boston: Digital Press.
- Grüning B, Chilton J, Köster J, Dale R, Soranzo N, Beek M van den, Goecks J, Backofen R, Nekrutenko A, Taylor J. 2018. Practical computational reproducibility in the life sciences. *Cell Systems* **6**:631–635. doi:[10.1016/j.cels.2018.03.014](#)
- Haroz S. 2018. Open Practices in Visualization Research (preprint). Open Science Framework. doi:[10.31219/osf.io/8ag3w](#)
- Holmes DT, Mobini M, McCudden CR. 2021. Reproducible manuscript preparation with RMarkdown application to JMSACL and other Elsevier Journals. *Journal of Mass Spectrometry and Advances in the Clinical Lab* **22**:8–16. doi:[10.1016/j.jmsacl.2021.09.002](#)
- Houtkoop BL, Chambers C, Macleod M, Bishop DVM, Nichols TE, Wagenmakers E-J. 2018. Data Sharing in Psychology: A Survey on Barriers and Preconditions. *Advances in Methods and Practices in Psychological Science* **1**:70–85. doi:[10.1177/2515245917751886](#)
- Ioannidis JPA. 2005. Why Most Published Research Findings Are False. *PLoS Medicine* **2**:e124. doi:[10.1371/journal.pmed.0020124](#)
- Jiménez RC, Kuzak M, Alhamdoosh M, Barker M, Batut B, Borg M, Capella-Gutierrez S, Chue Hong N, Cook M, Corpas M, Flannery M, Garcia L, Gelpí JLI, Gladman S, Goble C, González Ferreiro M, Gonzalez-Beltran A, Griffin PC, Grüning B, Hagberg J, Holub P, Hooft R, Ison J, Katz DS, Leskošek B, López Gómez F, Oliveira LJ, Mellor D, Mosbergen R, Mulder N, Perez-Riverol Y, Pergl R, Pichler H, Pope B, Sanz F, Schneider MV, Stodden V, Suchacki R, Svobodová Vařeková R, Talvik H-A, Todorov I, Treloar A, Tyagi S, Gompel M van, Vaughan D, Via A, Wang X, Watson-Haigh NS, Crouch S. 2017. Four simple recommendations to encourage best practices in research software. *F1000Research* **6**:876. doi:[10.12688/f1000research.11407.1](#)
- Judd CM, Westfall J, Kenny DA. 2017. Experiments with More Than One Random Factor: Designs, Analytic Models, and Statistical Power. *Annual Review of Psychology* **68**:601–625. doi:[10.1146/annurev-psych-122414-033702](#)



- Klein O, Hardwicke TE, Aust F, Breuer J, Danielsson H, Mohr AH, IJzerman H, Nilsson G, Vanpaemel W, Frank MC. 2018. A Practical Guide for Transparency in Psychological Science. *Collabra: Psychology* **4**:20. doi:[10.1525/collabra.158](https://doi.org/10.1525/collabra.158)
- Knuth DE. 1984. Literate Programming. *The Computer Journal* **27**:97–111. doi:[10.1093/comjnl/27.2.97](https://doi.org/10.1093/comjnl/27.2.97)
- Kosara R, Haroz S. 2018. Skipping the Replication Crisis in Visualization: Threats to Study Validity and How to Address Them : Position Paper2018 IEEE Evaluation and Beyond - Methodological Approaches for Visualization (BELIV). Berlin, Germany: IEEE. pp. 102–107. doi:[10.1109/BELIV.2018.8634392](https://doi.org/10.1109/BELIV.2018.8634392)
- Lam H, Bertini E, Isenberg P, Plaisant C, Carpendale S. 2012. Empirical Studies in Information Visualization: Seven Scenarios. *IEEE Transactions on Visualization and Computer Graphics* **18**:1520–1536. doi:[10.1109/TVCG.2011.279](https://doi.org/10.1109/TVCG.2011.279)
- Lamprecht A-L, Garcia L, Kuzak M, Martinez C, Arcila R, Martin Del Pico E, Dominguez Del Angel V, Van De Sandt S, Ison J, Martinez PA, McQuilton P, Valencia A, Harrow J, Psomopoulos F, Gelpi JLI, Chue Hong N, Goble C, Capella-Gutierrez S. 2020. Towards FAIR principles for research software. *Data Science* **3**:37–59. doi:[10.3233/DS-190026](https://doi.org/10.3233/DS-190026)
- Lee BD. 2018. Ten simple rules for documenting scientific software. *PLOS Computational Biology* **14**:e1006561. doi:[10.1371/journal.pcbi.1006561](https://doi.org/10.1371/journal.pcbi.1006561)
- Leipzig J, Nüst D, Hoyt CT, Ram K, Greenberg J. 2021. The role of metadata in reproducible computational research. *Patterns (New York, NY)* **2**:100322. doi:[10.1016/j.patter.2021.100322](https://doi.org/10.1016/j.patter.2021.100322)
- Lissa CJ van, Brandmaier AM, Brinkman L, Lamprecht A-L, Peikert A, Struiksma M, Vreede B. 2020. WORCS: A Workflow for Open Reproducible Code in Science. doi:[10.31234/osf.io/k4wde](https://doi.org/10.31234/osf.io/k4wde)
- Marwick B, Boettiger C, Mullen L. 2018. Packaging Data Analytical Work Reproducibly Using R (and Friends). *The American Statistician* **72**:80–88. doi:[10.1080/00031305.2017.1375986](https://doi.org/10.1080/00031305.2017.1375986)
- Mayrhofer R, Kuhbandner C, Lindner C. 2021. The Practice of Experimental Psychology: An Inevitably Postmodern Endeavor. *Frontiers in Psychology* **11**:612805. doi:[10.3389/fpsyg.2020.612805](https://doi.org/10.3389/fpsyg.2020.612805)
- Merkel D. 2014. Docker: Lightweight Linux Containers for Consistent Development and Deployment. *Linux J* **2014**.
- Meteyard L, Davies RAI. 2020. Best practice guidance for linear mixed-effects

- models in psychological science. *Journal of Memory and Language* **112**:104092. doi:[10.1016/j.jml.2020.104092](https://doi.org/10.1016/j.jml.2020.104092)
- Munafò MR, Nosek BA, Bishop DVM, Button KS, Chambers CD, Percie Du Sert N, Simonsohn U, Wagenmakers E-J, Ware JJ, Ioannidis JPA. 2017. A manifesto for reproducible science. *Nature Human Behaviour* **1**:0021. doi:[10.1038/s41562-016-0021](https://doi.org/10.1038/s41562-016-0021)
- Nüst D, Eddelbuettel D, Bennett D, Cannoodt R, Clark D, Daroczi G, Edmondson M, Fay C, Hughes E, Kjeldgaard L, Lopp S, Marwick B, Nolis H, Nolis J, Ooi H, Ram K, Ross N, Shepherd L, Sólymos P, Swetnam TL, Turaga N, Van Petegem C, Williams J, Willis C, Xiao N. 2020a. The Rockerverse: Packages and Applications for Containerization with R. *The R Journal* **12**:437. doi:[10.32614/RJ-2020-007](https://doi.org/10.32614/RJ-2020-007)
- Nüst D, Konkol M, Pebesma E, Kray C, Schutzeichel M, Przibytzin H, Lorenz J. 2017. Opening the Publication Process with Executable Research Compendia. *D-Lib Magazine* **23**. doi:[10.1045/january2017-nuest](https://doi.org/10.1045/january2017-nuest)
- Nüst D, Sochat V, Marwick B, Eglen SJ, Head T, Hirst T, Evans BD. 2020b. Ten simple rules for writing Dockerfiles for reproducible data science. *PLOS Computational Biology* **16**:e1008316. doi:[10.1371/journal.pcbi.1008316](https://doi.org/10.1371/journal.pcbi.1008316)
- Open Science Collaboration. 2015. Estimating the reproducibility of psychological science. *Science* **349**:aac4716. doi:[10.1126/science.aac4716](https://doi.org/10.1126/science.aac4716)
- Peikert A, Brandmaier AM. 2021. A Reproducible Data Analysis Workflow. *Quantitative and Computational Methods in Behavioral Sciences* **1**:e3763. doi:[10.5964/qcmb.3763](https://doi.org/10.5964/qcmb.3763)
- Peirce J, Gray JR, Simpson S, MacAskill M, Höchenberger R, Sogo H, Kastman E, Lindeløv JK. 2019. PsychoPy2: Experiments in behavior made easy. *Behavior Research Methods* **51**:195–203. doi:[10.3758/s13428-018-01193-y](https://doi.org/10.3758/s13428-018-01193-y)
- Peng RD. 2011. Reproducible research in computational science. *Science* **334**:1226–1227. doi:[10.1126/science.1213847](https://doi.org/10.1126/science.1213847)
- Piccolo SR, Frampton MB. 2016. Tools and techniques for computational reproducibility. *GigaScience* **5**:30. doi:[10.1186/s13742-016-0135-4](https://doi.org/10.1186/s13742-016-0135-4)
- Plesser HE. 2018. Reproducibility vs. Replicability: A Brief History of a Confused Terminology. *Frontiers in Neuroinformatics* **11**:76. doi:[10.3389/fninf.2017.00076](https://doi.org/10.3389/fninf.2017.00076)
- Rule A, Birmingham A, Zuniga C, Altintas I, Huang S-C, Knight R, Moshiri N, Nguyen MH, Rosenthal SB, Pérez F, Rose PW. 2019. Ten simple rules for writing and sharing computational analyses in Jupyter Notebooks. *PLOS Computational Biology*

- 15:e1007007. doi:[10.1371/journal.pcbi.1007007](https://doi.org/10.1371/journal.pcbi.1007007)
- Sandve GK, Nekrutenko A, Taylor J, Hovig E. 2013. Ten Simple Rules for Reproducible Computational Research. *PLoS Computational Biology* **9**:e1003285. doi:[10.1371/journal.pcbi.1003285](https://doi.org/10.1371/journal.pcbi.1003285)
- Simmons JP, Nelson LD, Simonsohn U. 2011. False-Positive Psychology: Undisclosed Flexibility in Data Collection and Analysis Allows Presenting Anything as Significant. *Psychological Science* **22**:1359–1366. doi:[10.1177/0956797611417632](https://doi.org/10.1177/0956797611417632)
- Singmann H, Kellen D. 2019. An Introduction to Mixed Models for Experimental Psychology In: Spieler D, Schumacher E, editors. *New Methods in Cognitive Psychology*. Routledge. pp. 4–31. doi:[10.4324/9780429318405-2](https://doi.org/10.4324/9780429318405-2)
- Stodden V, McNutt M, Bailey DH, Deelman E, Gil Y, Hanson B, Heroux MA, Ioannidis JPA, Taufer M. 2016. Enhancing reproducibility for computational methods. *Science* **354**:1240–1241. doi:[10.1126/science.aah6168](https://doi.org/10.1126/science.aah6168)
- Stodden V, Seiler J, Ma Z. 2018. An empirical analysis of journal policy effectiveness for computational reproducibility. *Proceedings of the National Academy of Sciences* **115**:2584–2589. doi:[10.1073/pnas.1708290115](https://doi.org/10.1073/pnas.1708290115)
- Trisovic A, Lau MK, Pasquier T, Crosas M. 2022. A large-scale study on research code quality and execution. *Scientific Data* **9**:60. doi:[10.1038/s41597-022-01143-6](https://doi.org/10.1038/s41597-022-01143-6)
- Trisovic A, Mika K, Boyd C, Feger S, Crosas M. 2021. Repository Approaches to Improving the Quality of Shared Data and Code. *Data* **6**:15. doi:[10.3390/data6020015](https://doi.org/10.3390/data6020015)
- Ushey K, Wickham H. 2023. [Renv: Project Environments](#).
- Voeten CC. 2022. [Buildmer: Stepwise Elimination and Term Reordering for Mixed-Effects](#).
- White E, Baldridge E, Brym Z, Locey K, McGlinn D, Supp S. 2013. Nine simple ways to make it easier to (re)use your data. *Ideas in Ecology and Evolution* **6**. doi:[10.4033/iee.2013.6b.6.f](https://doi.org/10.4033/iee.2013.6b.6.f)
- Wickham H. 2014. Tidy Data. *Journal of Statistical Software* **59**. doi:[10.18637/jss.v059.i10](https://doi.org/10.18637/jss.v059.i10)
- Wickham H, Averick M, Bryan J, Chang W, McGowan L, François R, Golemund G, Hayes A, Henry L, Hester J, Kuhn M, Pedersen T, Miller E, Bache S, Müller K, Ooms J, Robinson D, Seidel D, Spinu V, Takahashi K, Vaughan D, Wilke C, Woo K, Yutani H. 2019. Welcome to the Tidyverse. *Journal of Open Source Software* **4**:1686. doi:[10.21105/joss.01686](https://doi.org/10.21105/joss.01686)

- Wilkinson L, Task Force on Statistical Inference. 1999. Statistical methods in psychology journals: Guidelines and explanations. *American Psychologist* **54**:594–604. doi:[10.1037/0003-066X.54.8.594](https://doi.org/10.1037/0003-066X.54.8.594)
- Wilkinson MD, Dumontier M, Aalbersberg IJ, Appleton G, Axton M, Baak A, Blomberg N, Boiten J-W, Da Silva Santos LB, Bourne PE, Bouwman J, Brookes AJ, Clark T, Crosas M, Dillo I, Dumon O, Edmunds S, Evelo CT, Finkers R, Gonzalez-Beltran A, Gray AJG, Groth P, Goble C, Grethe JS, Heringa J, Hoen PAC 'T, Hooft R, Kuhn T, Kok R, Kok J, Lusher SJ, Martone ME, Mons A, Packer AL, Persson B, Rocca-Serra P, Roos M, Van Schaik R, Sansone S-A, Schultes E, Sengstag T, Slater T, Strawn G, Swertz MA, Thompson M, Van Der Lei J, Van Mulligen E, Velterop J, Waagmeester A, Wittenburg P, Wolstencroft K, Zhao J, Mons B. 2016. The FAIR Guiding Principles for scientific data management and stewardship. *Scientific Data* **3**:160018. doi:[10.1038/sdata.2016.18](https://doi.org/10.1038/sdata.2016.18)
- Wilson G, Bryan J, Cranston K, Kitzes J, Nederbragt L, Teal TK. 2017. Good enough practices in scientific computing. *PLOS Computational Biology* **13**:e1005510. doi:[10.1371/journal.pcbi.1005510](https://doi.org/10.1371/journal.pcbi.1005510)

## Preface to Chapter 4

This thesis presents a series of empirical experiments which investigate the cognitive processing of the absolute magnitude of values presented in data visualisations. Conceptions of absolute magnitude refer to how *large* or *small* numerical values are. The first set of experiments in this thesis provides the foundation for this research project.

Prior research has reported that values presented at higher positions were judged as higher in magnitude than values presented at lower positions (Sandman et al., 1994). The first set of experiments in this thesis also explores the role of *physical position* in informing magnitude judgements. I attempt to replicate Sandman et al.'s (1994) general finding, then expand upon these results to examine the underlying cognitive mechanism.

In these experiments, participants observed a series of simple dot plot visualisations of fictitious data, with numerical values encoded through their position on the vertical axis. Systematically manipulating the axis limits surrounding plotted values causes the same values to appear at high or low positions. Therefore, participants' ratings of the magnitude of plotted values reveals the effect of this design choice on interpretations. The simplicity of the stimuli provides high experimental control, but the use of meaningful scenarios concerning risk-related events increases ecological validity.

In order to gain further insight into how absolute magnitude is processed, these experiments also examine whether the *physical* positions or *relative* positions of plotted values contribute to magnitude judgements. This distinction concerns whether interpretations are influenced primarily by the association between high physical positions and high values, or the numerical context within which values are plotted. Dot plots with *inverted* axes are presented to test these two competing explanations.

## References

Sandman PM, Weinstein ND, Miller P. 1994. High Risk or Low: How Location on a "Risk Ladder" Affects Perceived Risk. *Risk Analysis* **14**:35–45. doi:[10.1111/j.1539-6924.1994.tb00026.x](https://doi.org/10.1111/j.1539-6924.1994.tb00026.x)